

Artificiële Intelligentie

RCA AIG 04Q6 25 MAART 2021



*Computational Foundations of
Machine Learning [ML]
with Python*

CONTEXT

- **Prerequisites:** basics in linear algebra, probability, and analysis of algorithms.
- **Workload:** homework assignments
- **GitHub:** Start a ML repository at GitHub

This Lecture

- Basic definitions and concepts.
- Introduction to the problem of learning.

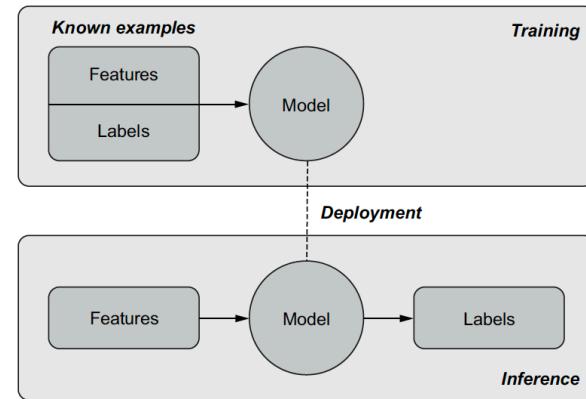
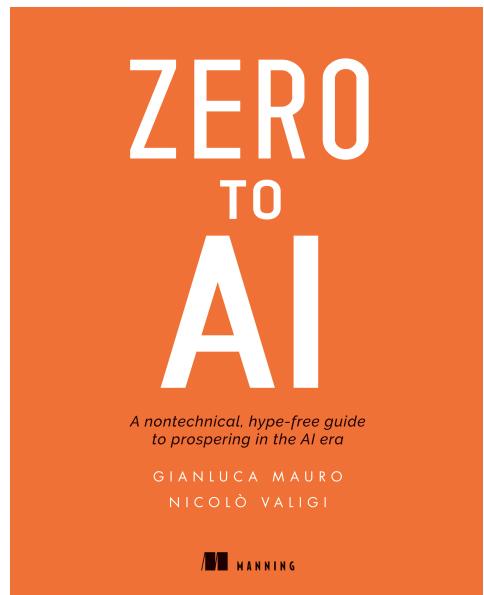


Figure 2.3 The two phases of machine learning: training and inference

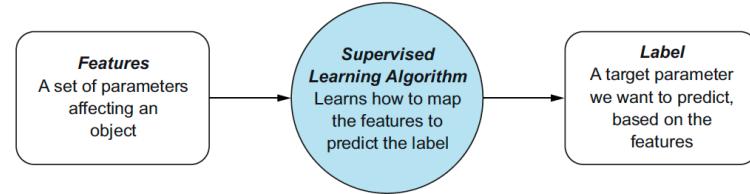


Figure 2.4 The core concept of supervised learning: finding a mapping between a set of features and a label

ZERO TO AI

A nontechnical, hype-free guide
to prospering in the AI era

GIANLUCA MAURO
NICOLÒ VALIGI

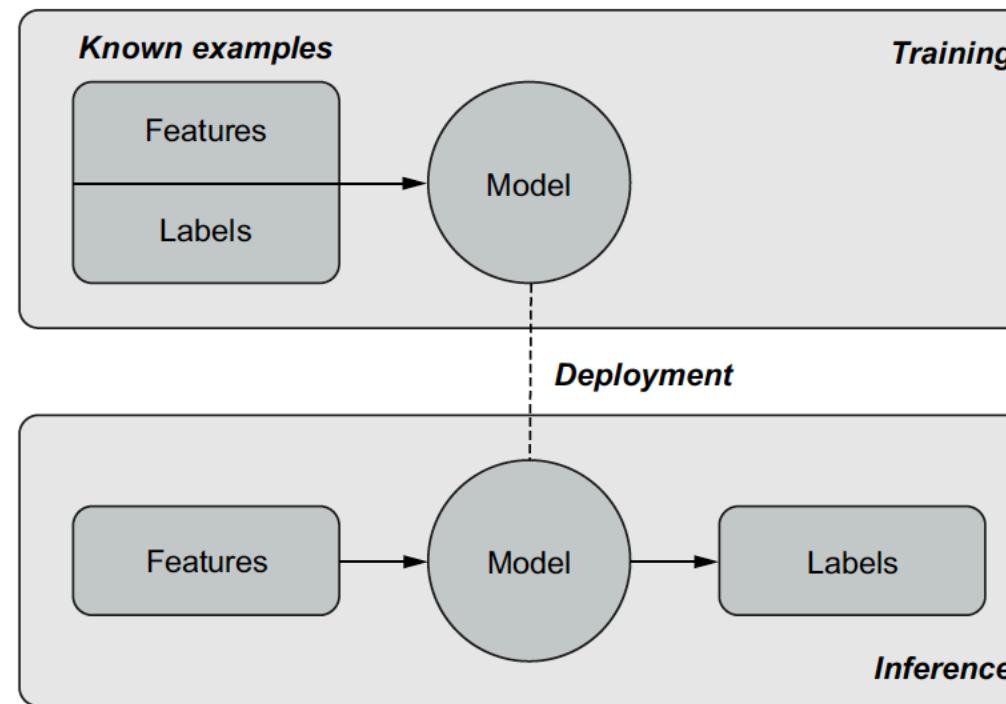


Figure 2.3 The two phases of machine learning: training and Inference

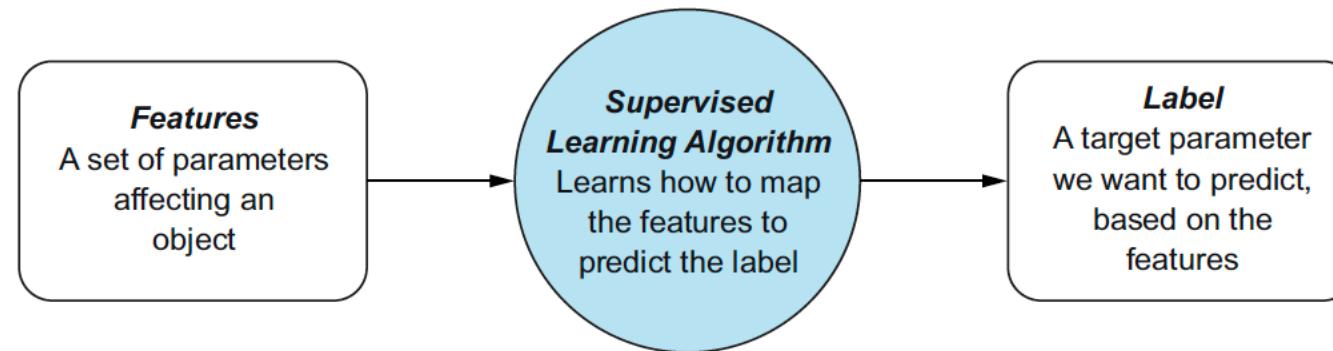
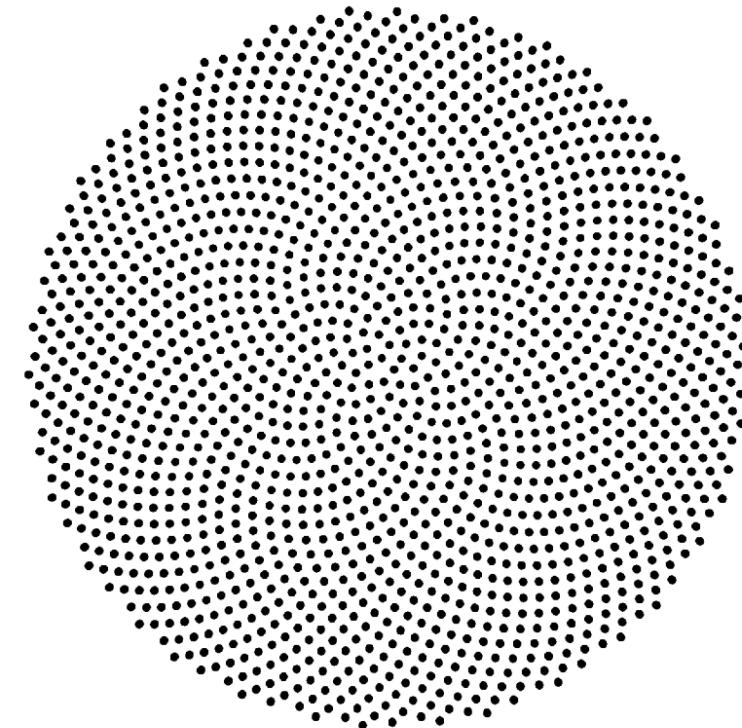
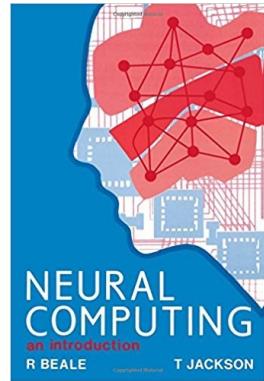
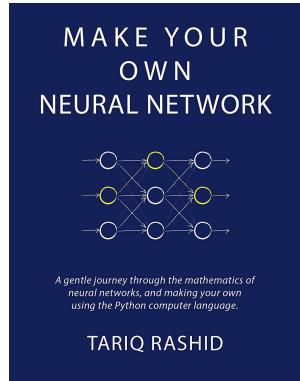


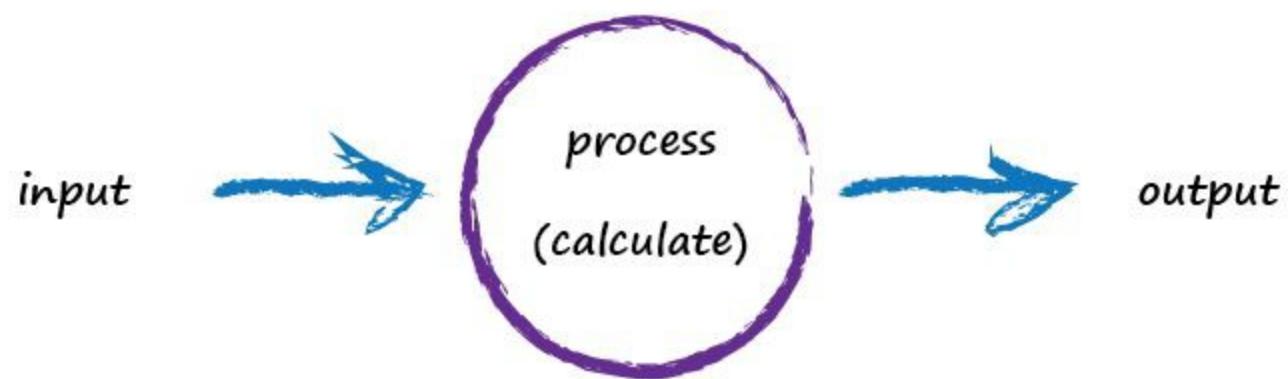
Figure 2.4 The core concept of supervised learning: finding a mapping between a set of features and a label

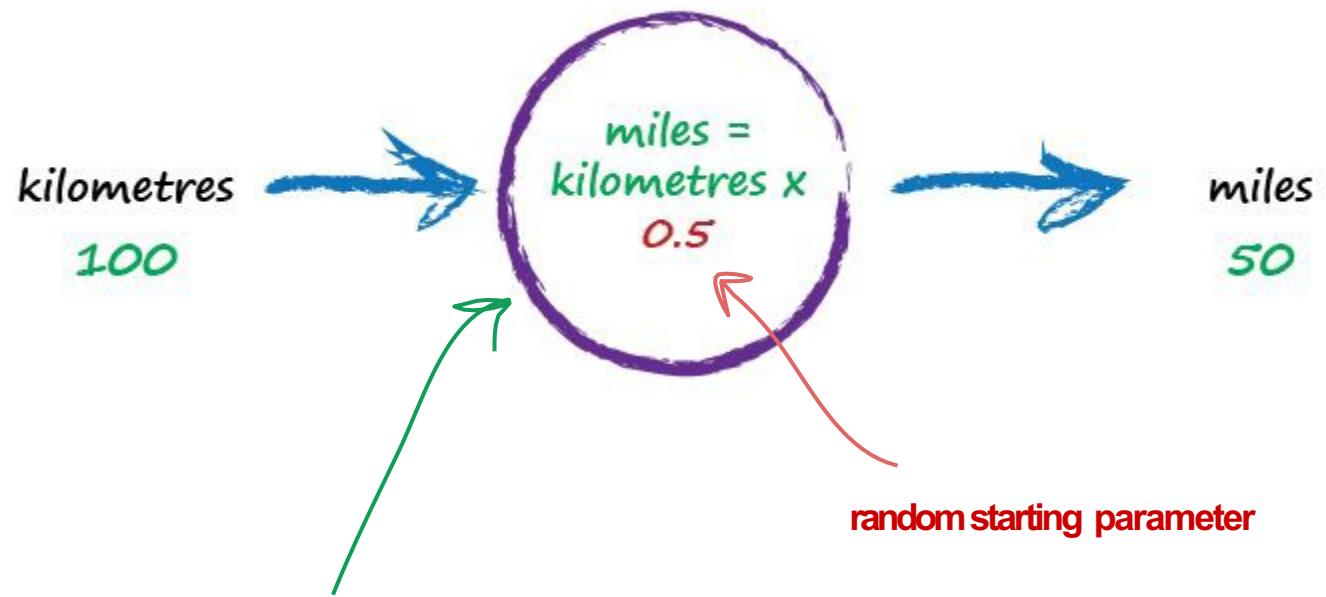
{01}

Fundamentals

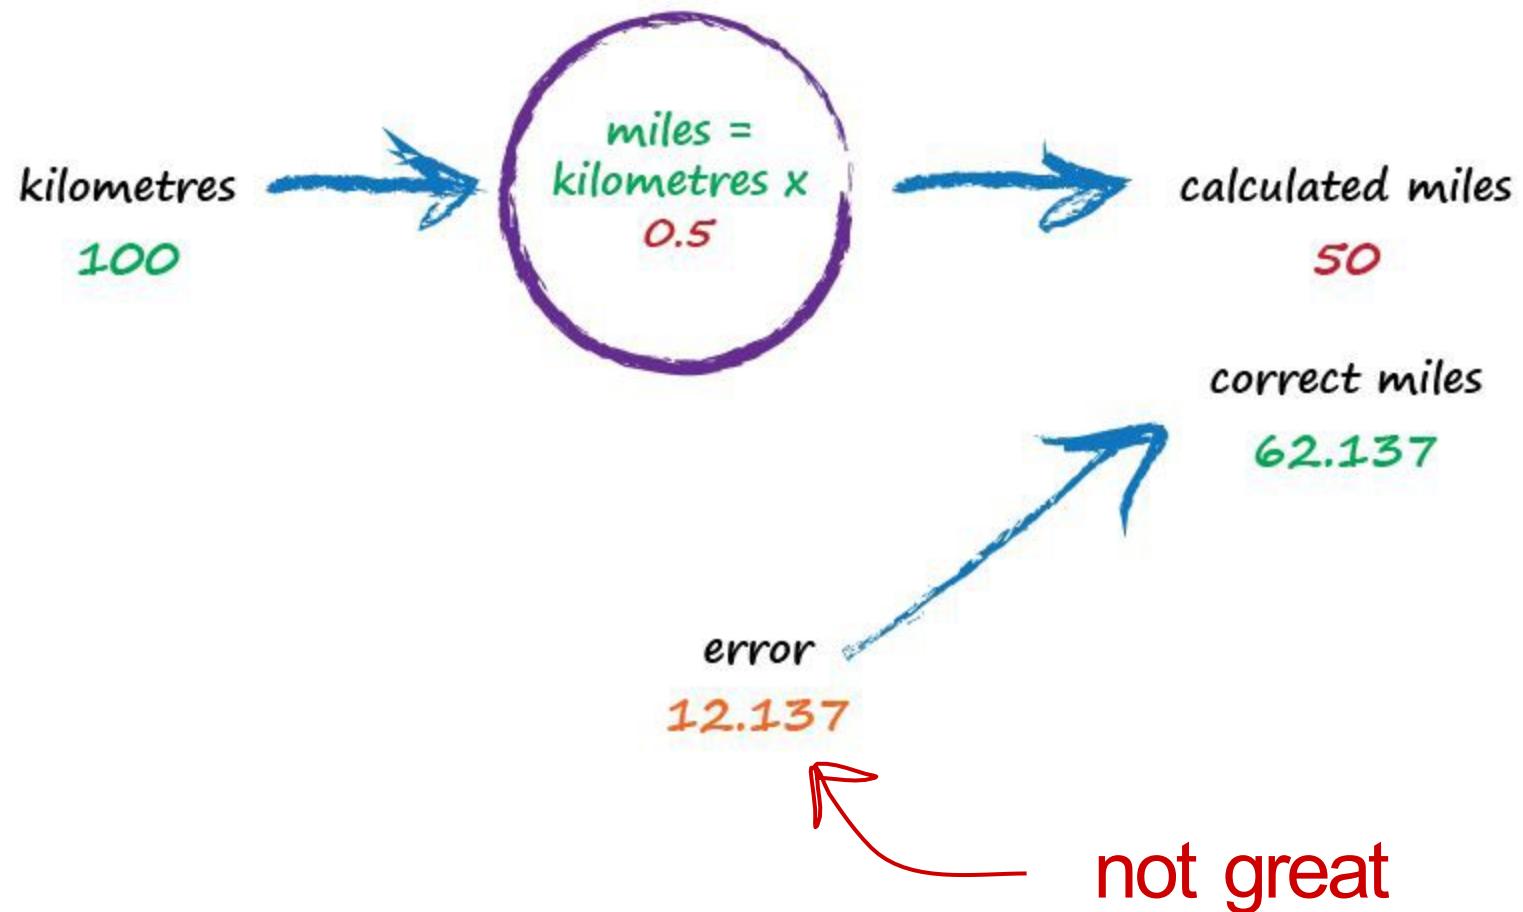


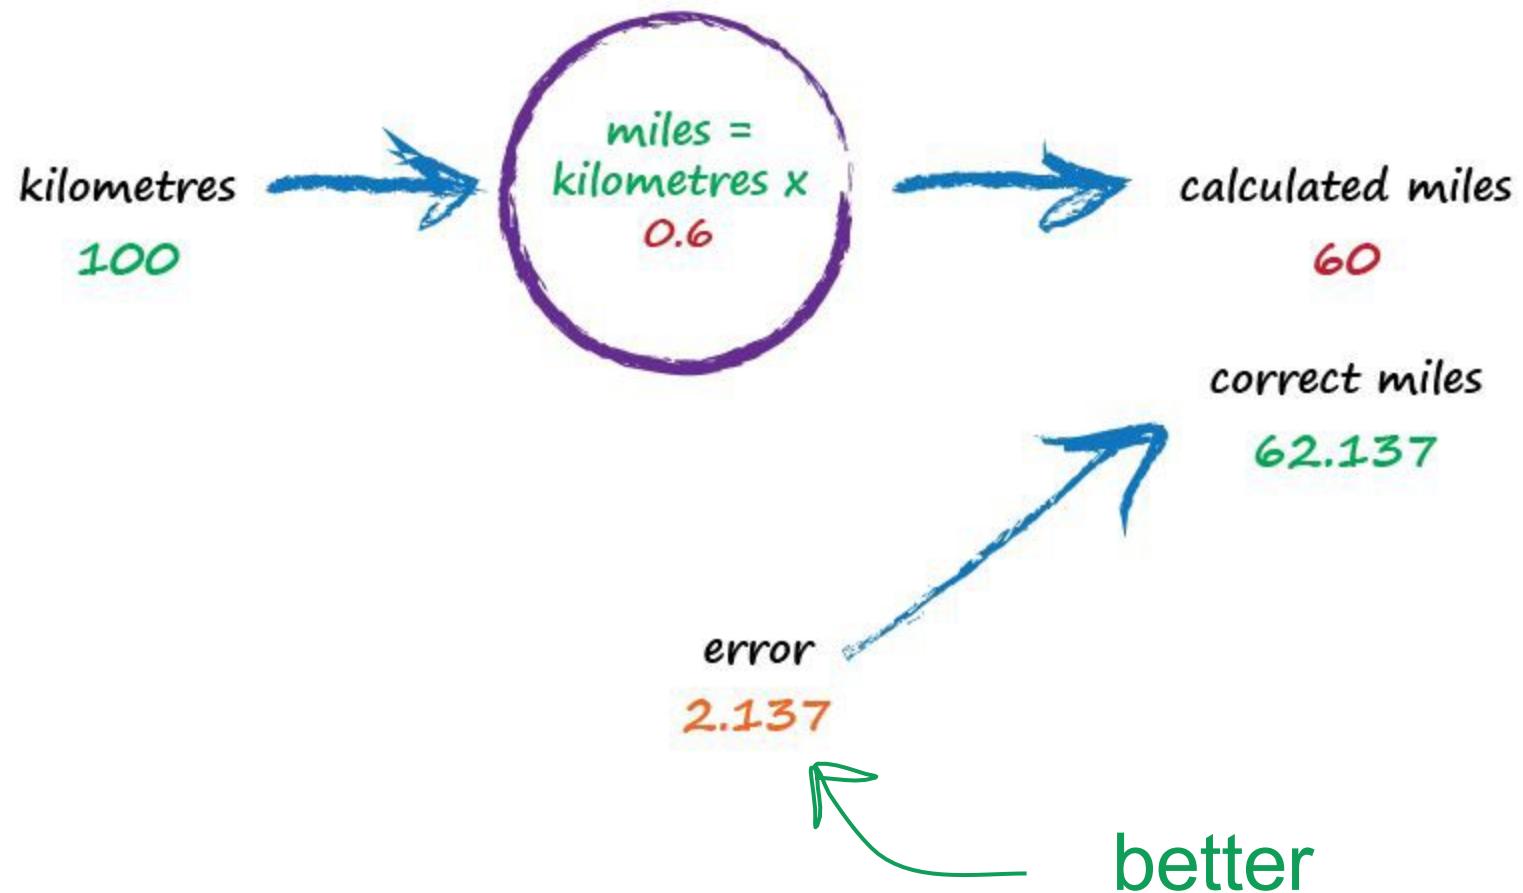


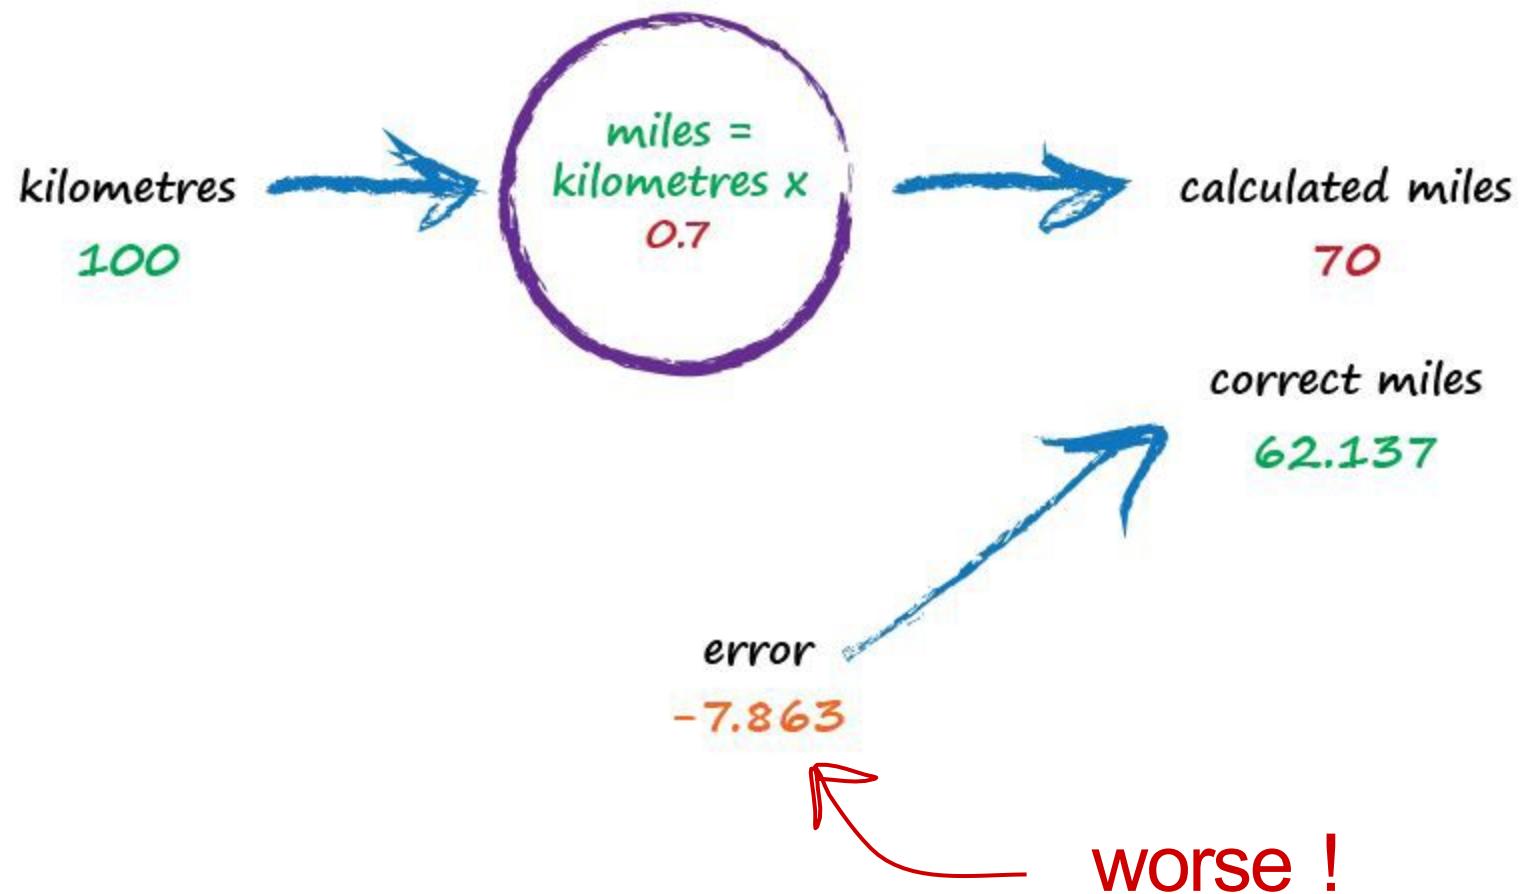


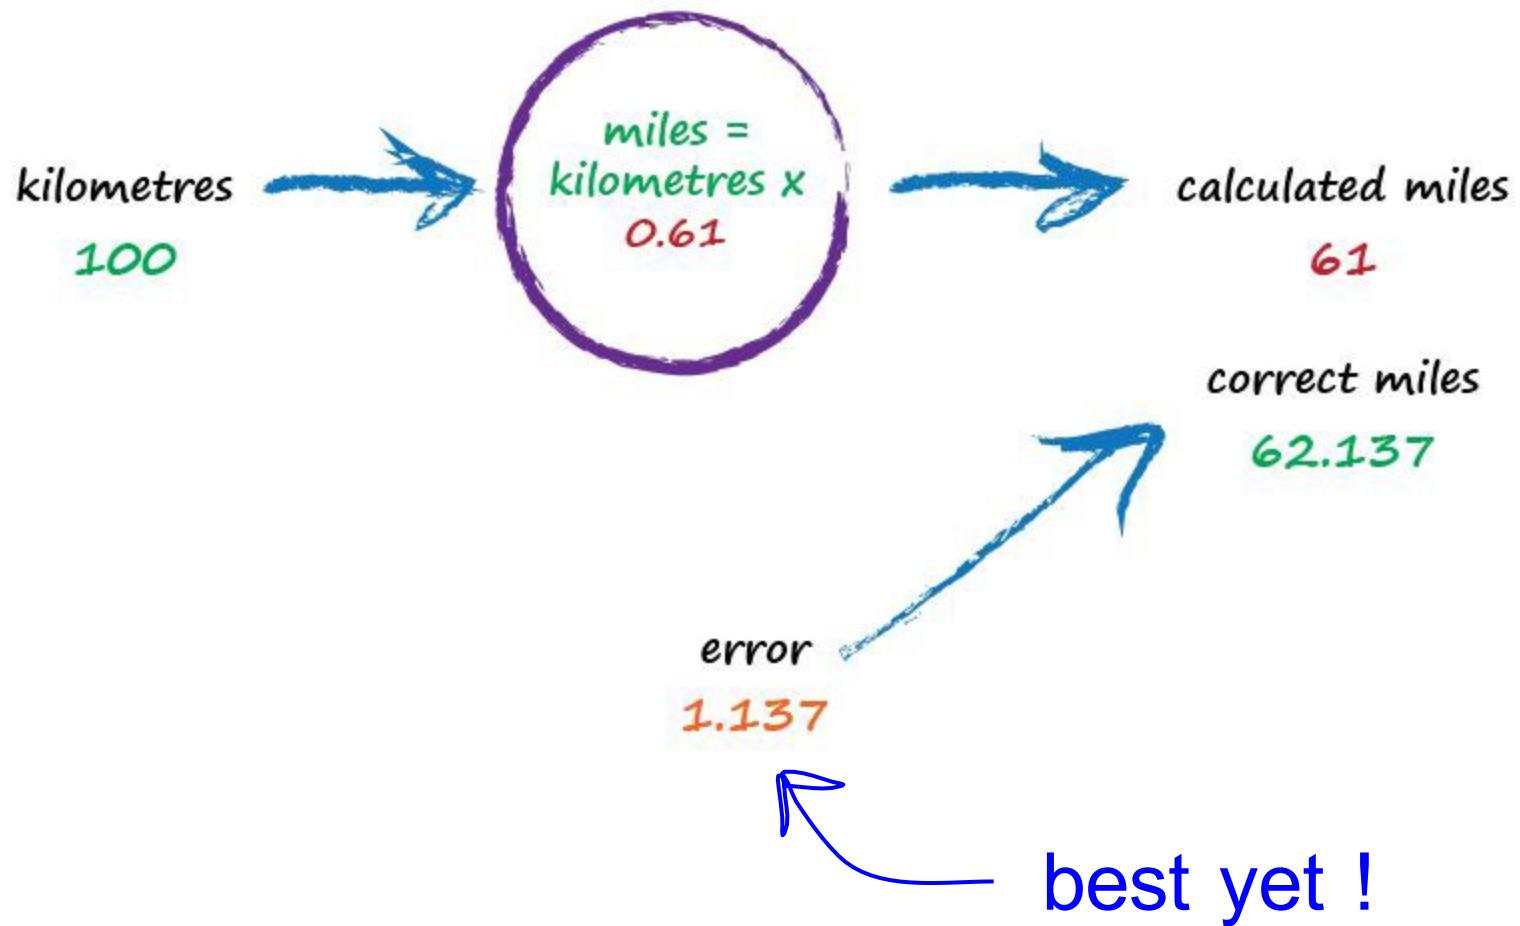


try a model - this one is linear





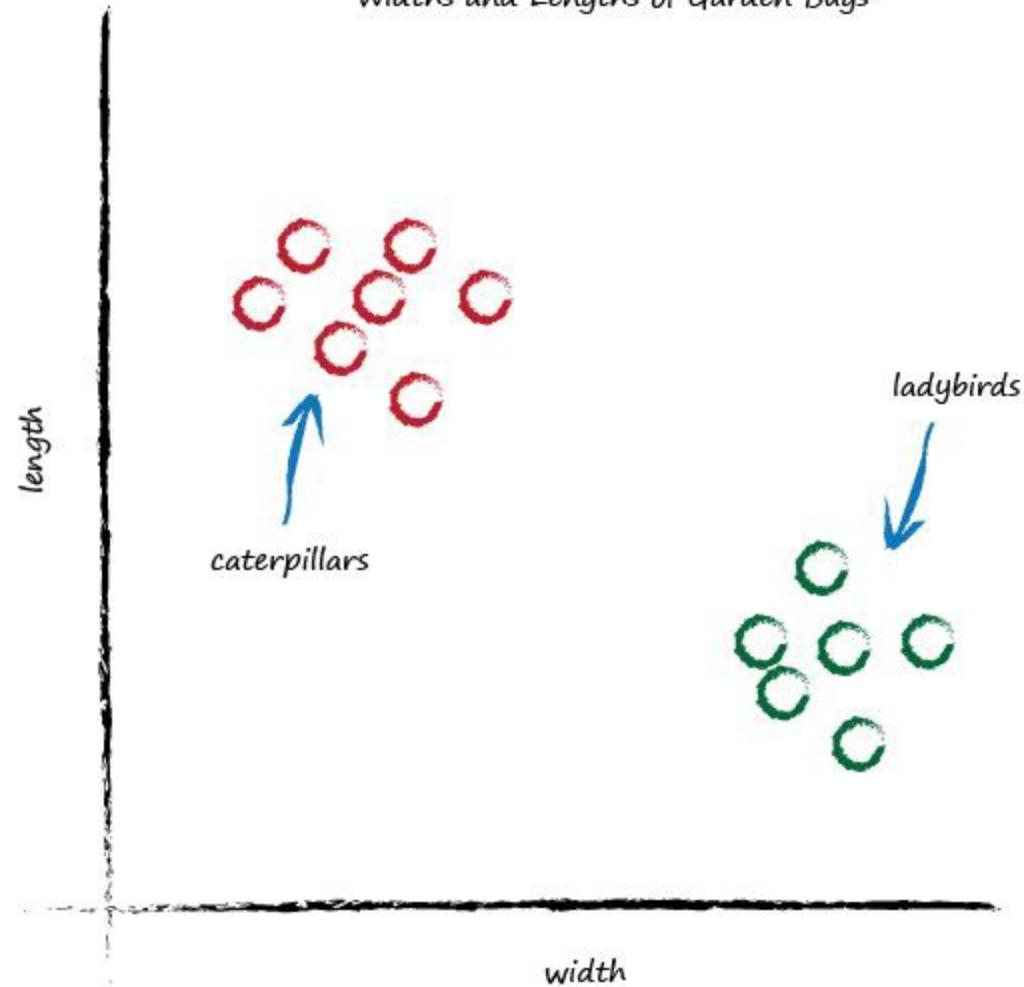




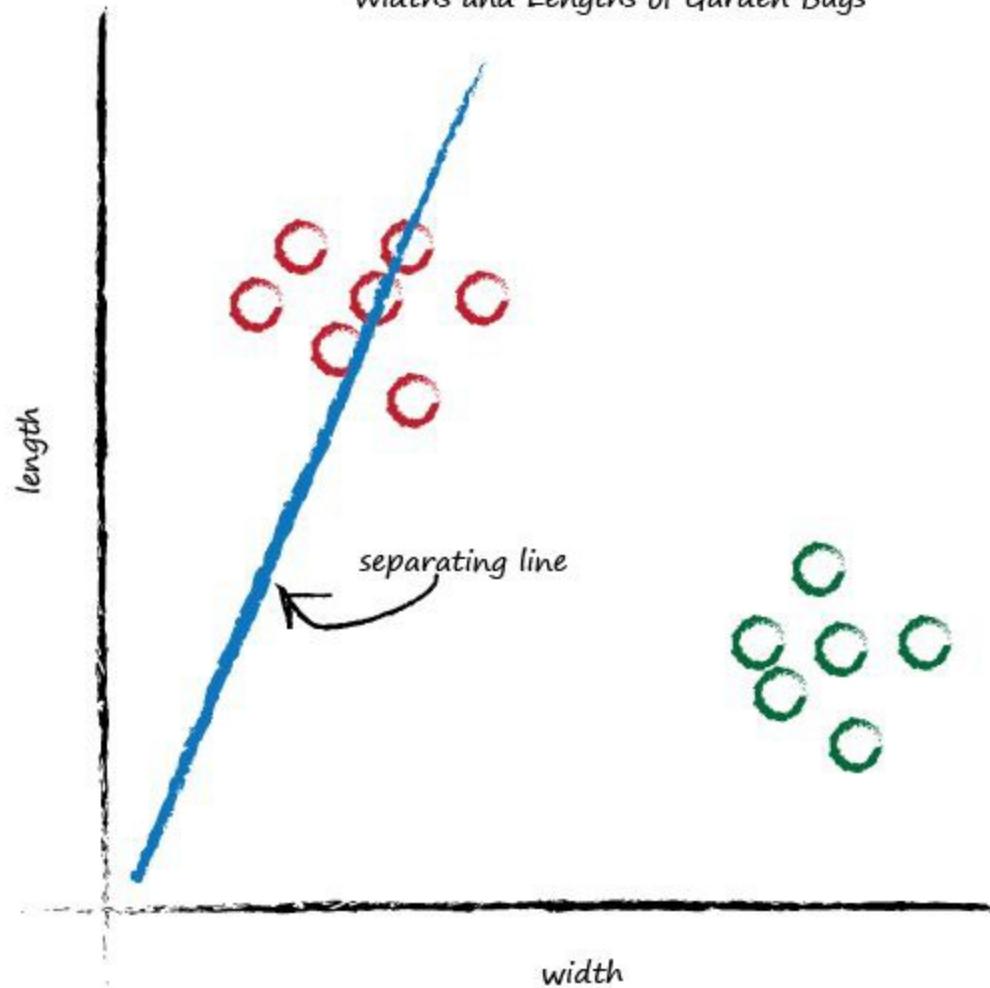


1. Don't know how something works exactly? Try a **model** with **adjustable** parameters.
2. Use the **error** to refine the parameters.

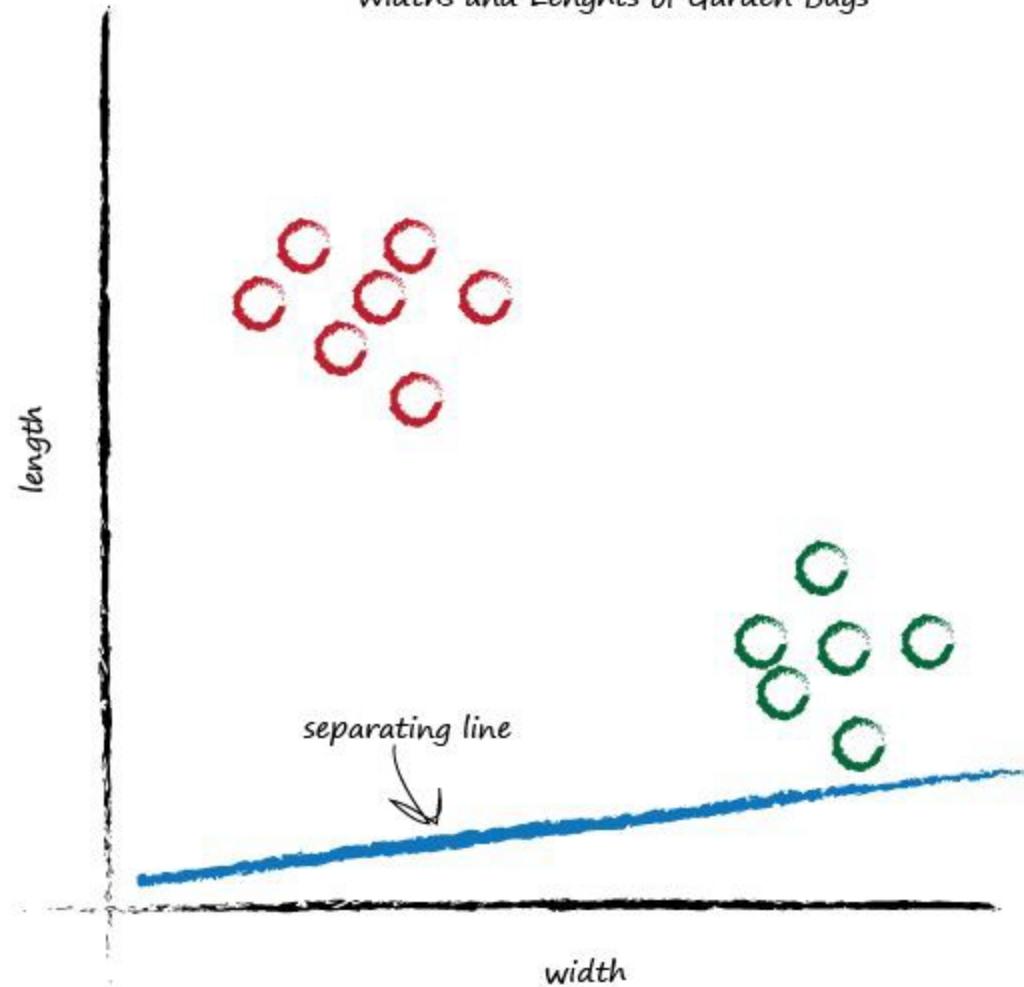
Widths and Lengths of Garden Bugs



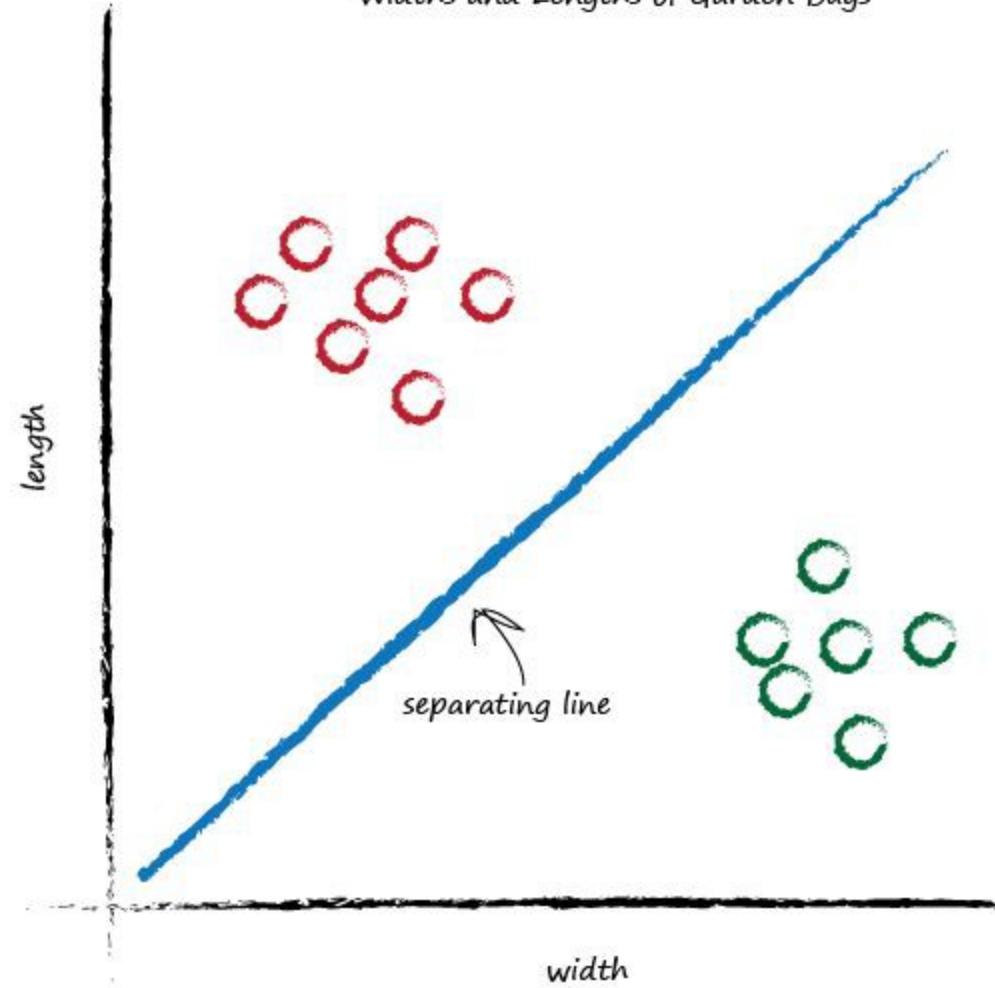
Widths and Lengths of Garden Bugs



Widths and Lengths of Garden Bugs



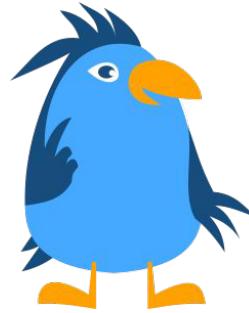
Widths and Lengths of Garden Bugs





1. **Classifying** things is kinda like **predicting** things.

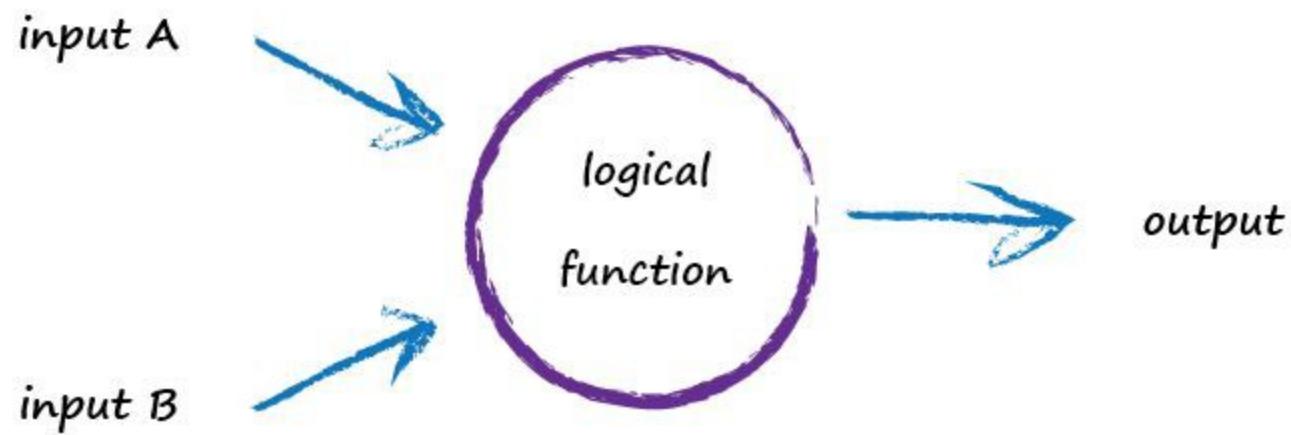
Logic

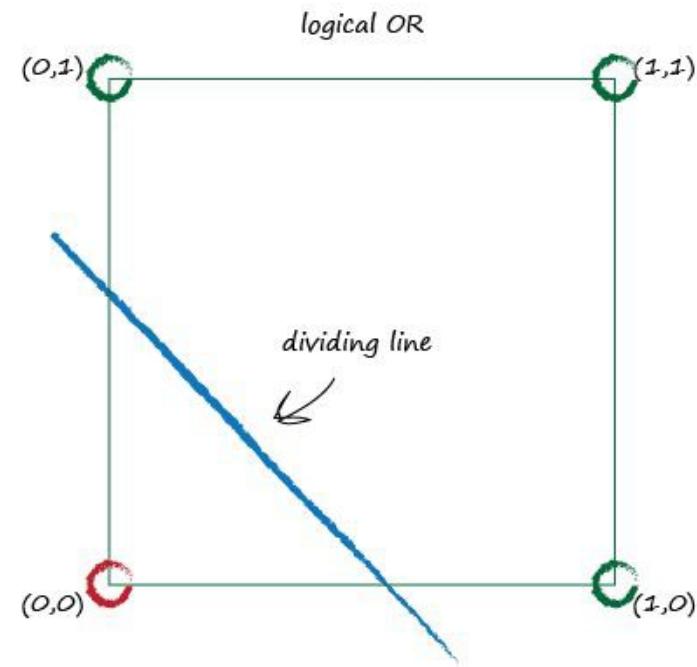
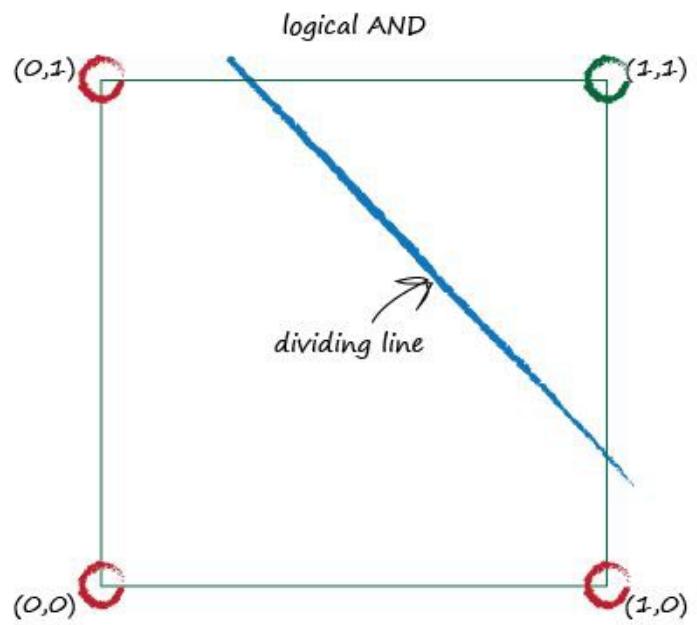


IF I have eaten my vegetables **AND** I am still hungry
THEN I can have ice cream.

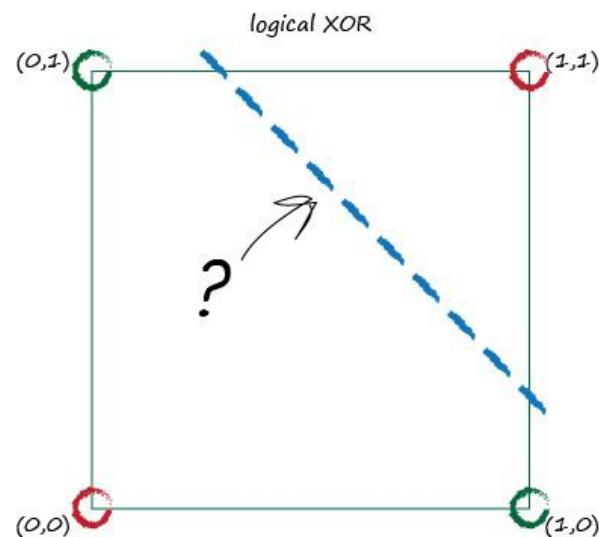
IF it's the weekend **OR** I am on annual leave **THEN** I'll go to the park.

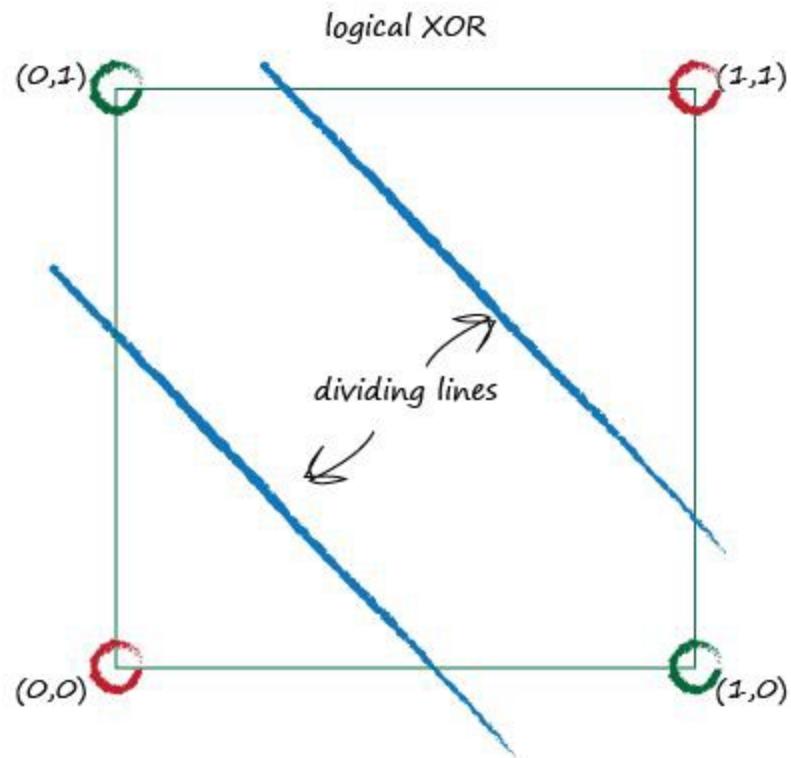
Input A	Input B	AND	OR
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	1





Input A	Input B	XOR
0	0	0
0	1	1
1	0	1
1	1	0



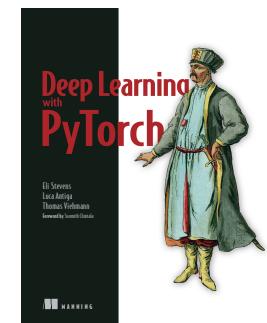
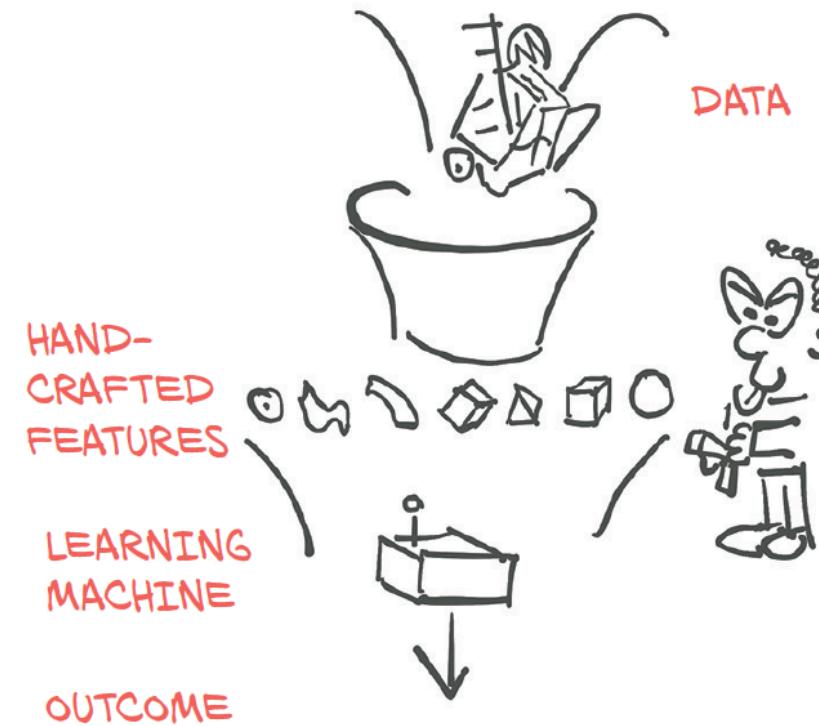


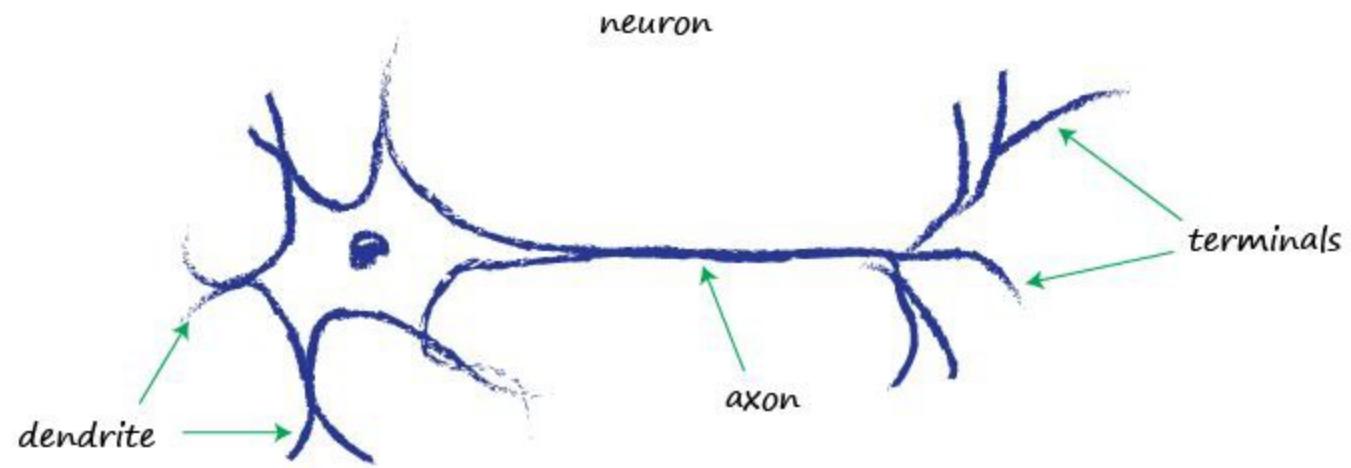
...Use more than one node!



1. Some problems can't be solved with just a single simple linear classifier.
2. You can use **multiple nodes** working together to solve many of these problems.

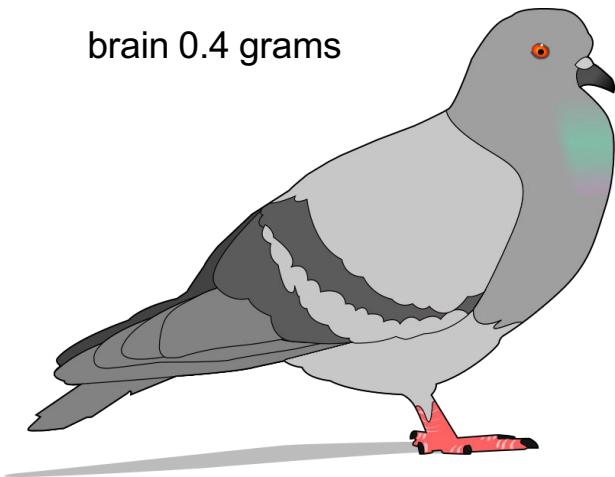
MACHINE LEARNING [ML]



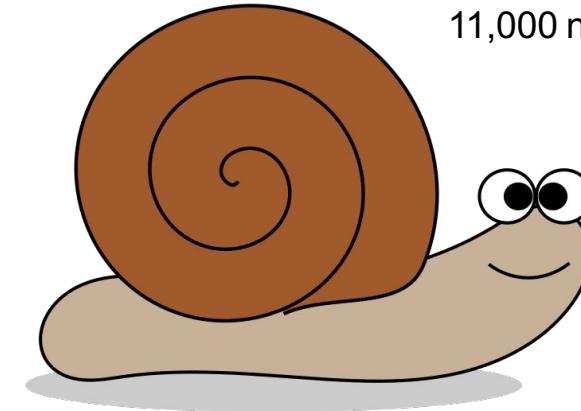


Nature

brain 0.4 grams



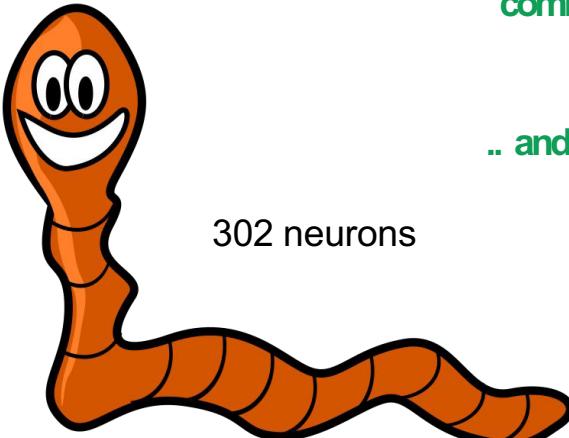
11,000 neurons



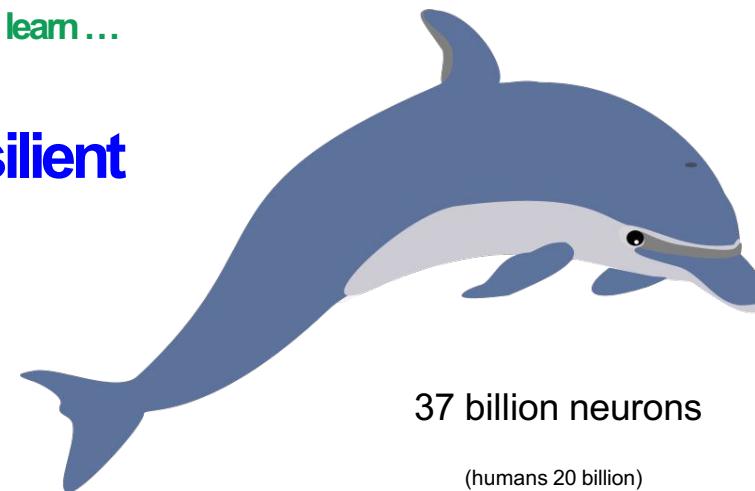
nature's brains can eat, fly, navigate, fight,
communicate, play, learn ...

.. and they're **resilient**

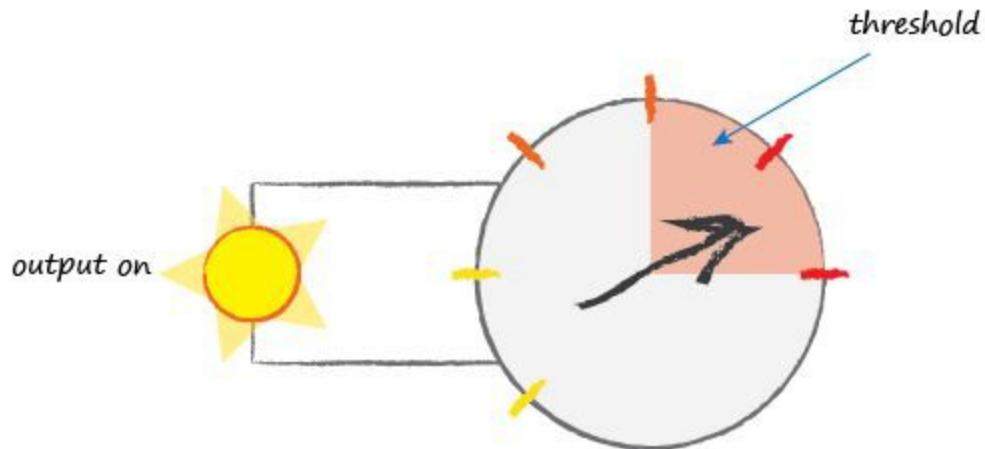
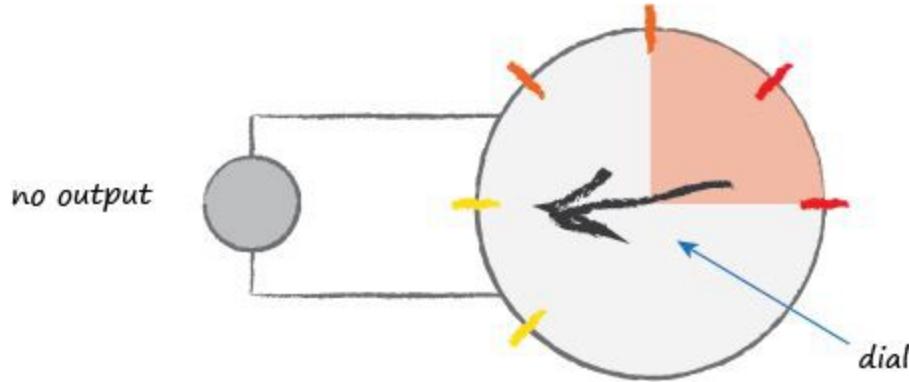
302 neurons

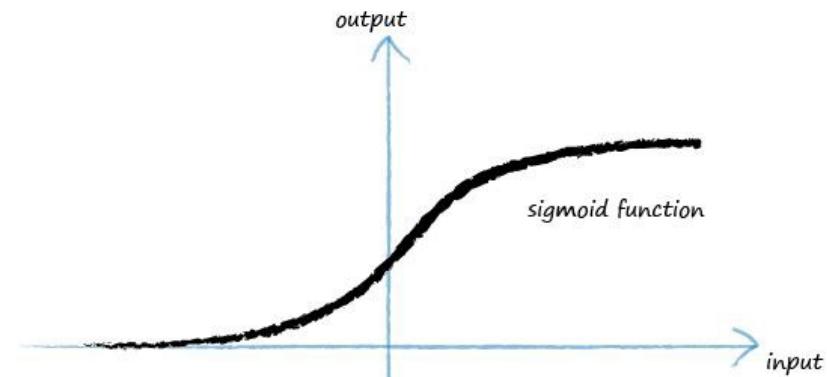
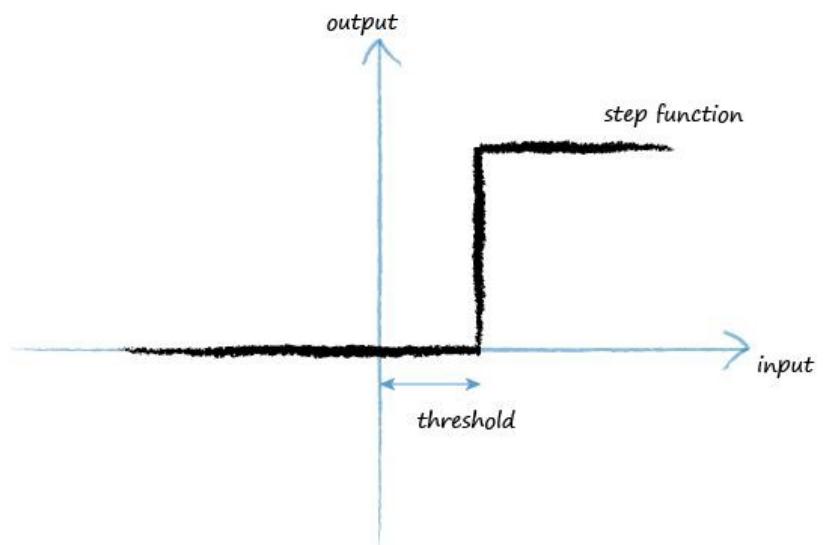


37 billion neurons



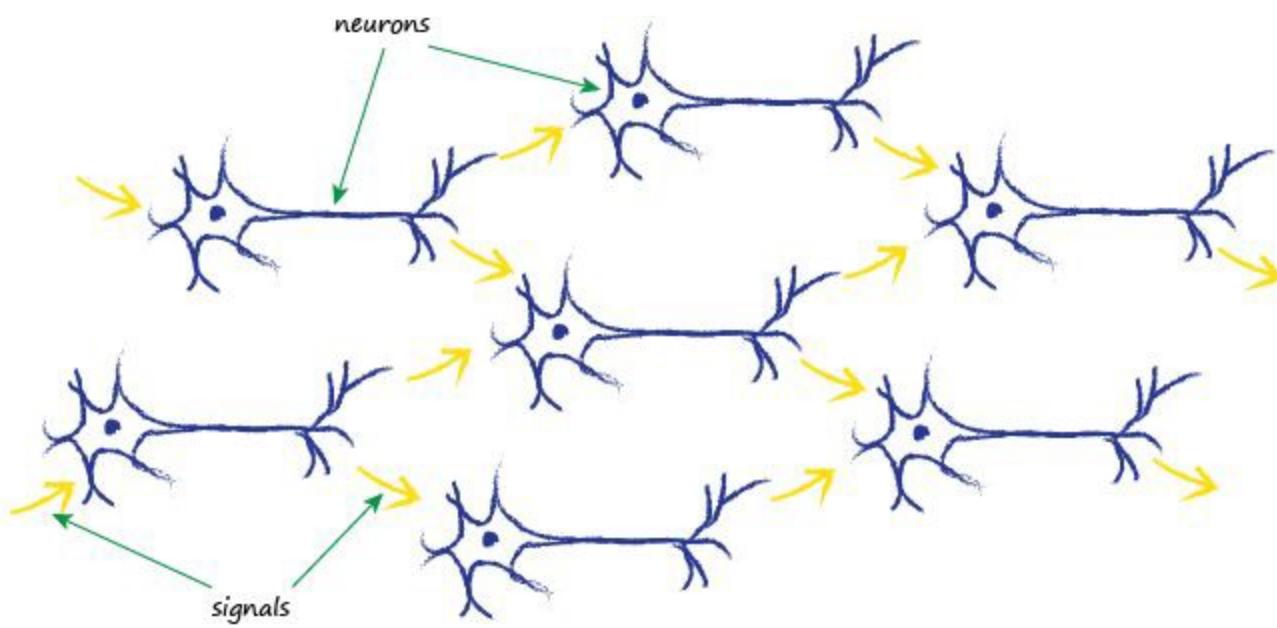
(humans 20 billion)

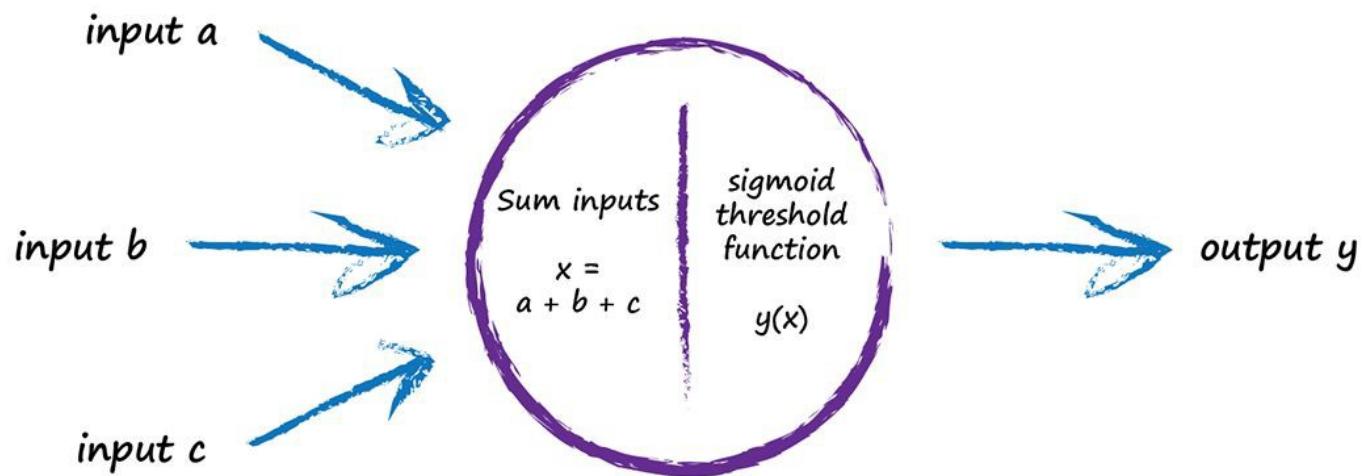


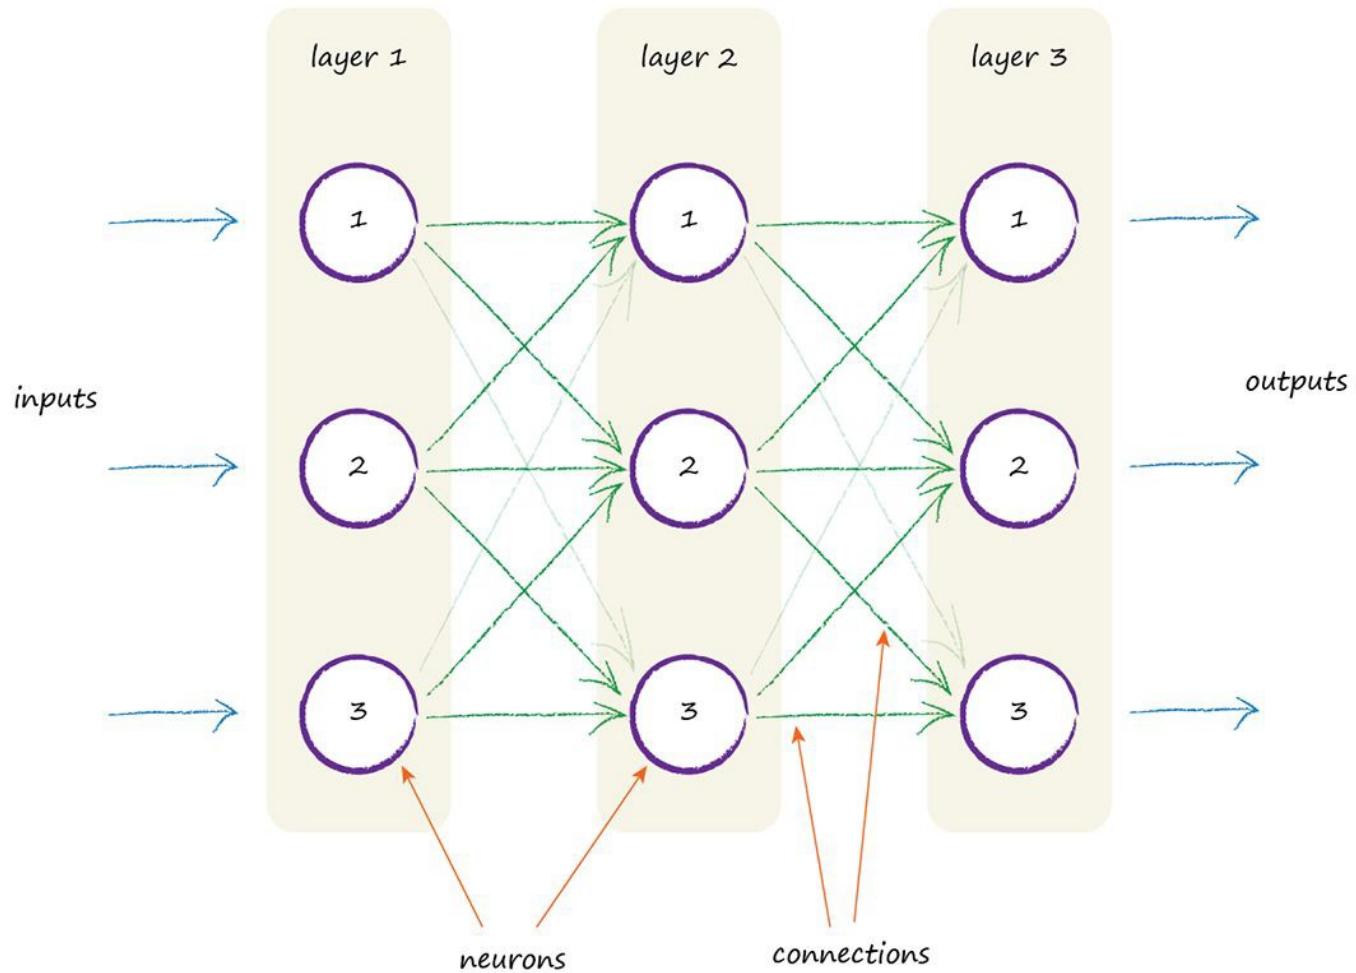


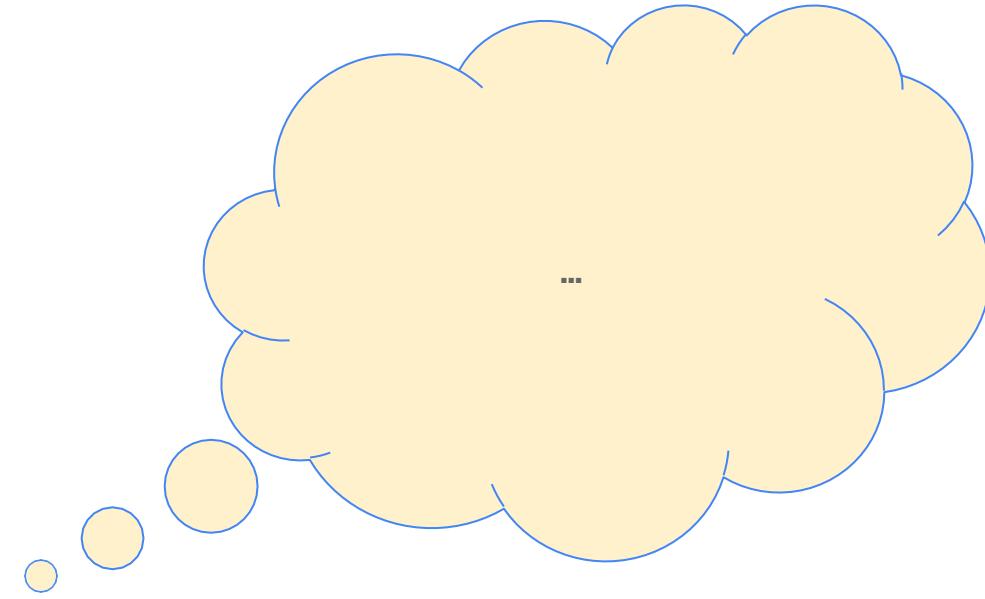
logistic function

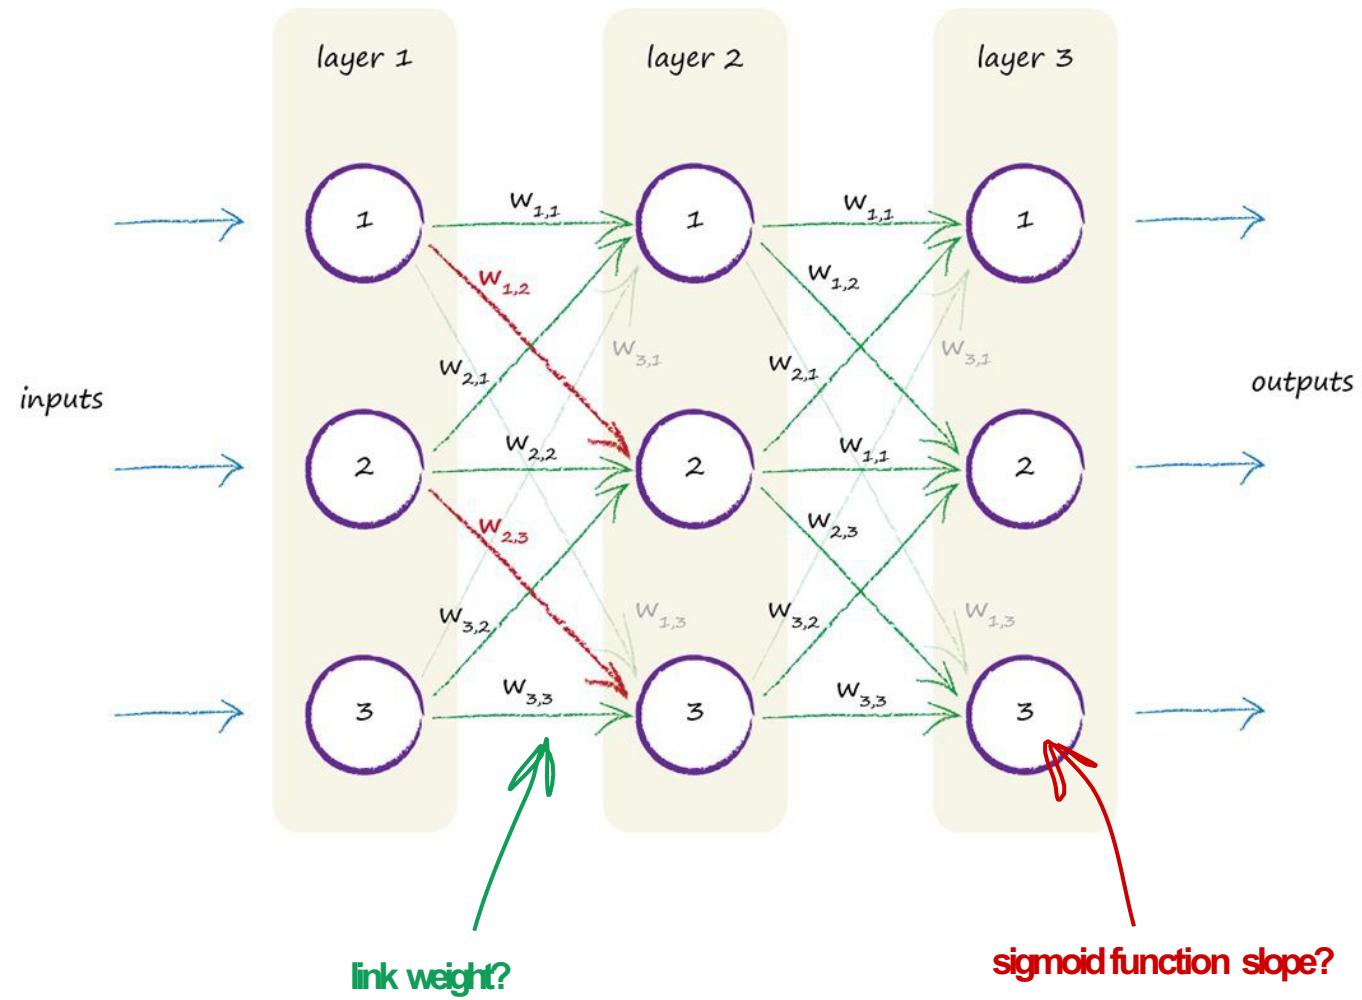
$$y = 1 / (1 + e^{-x})$$





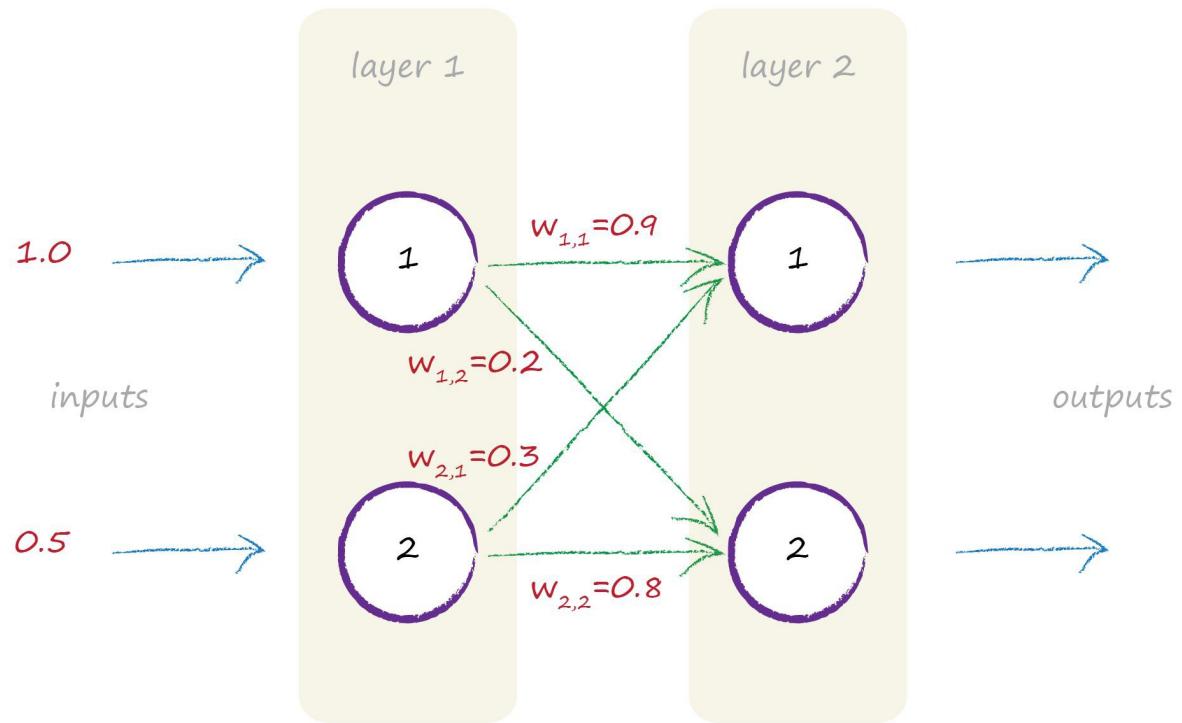


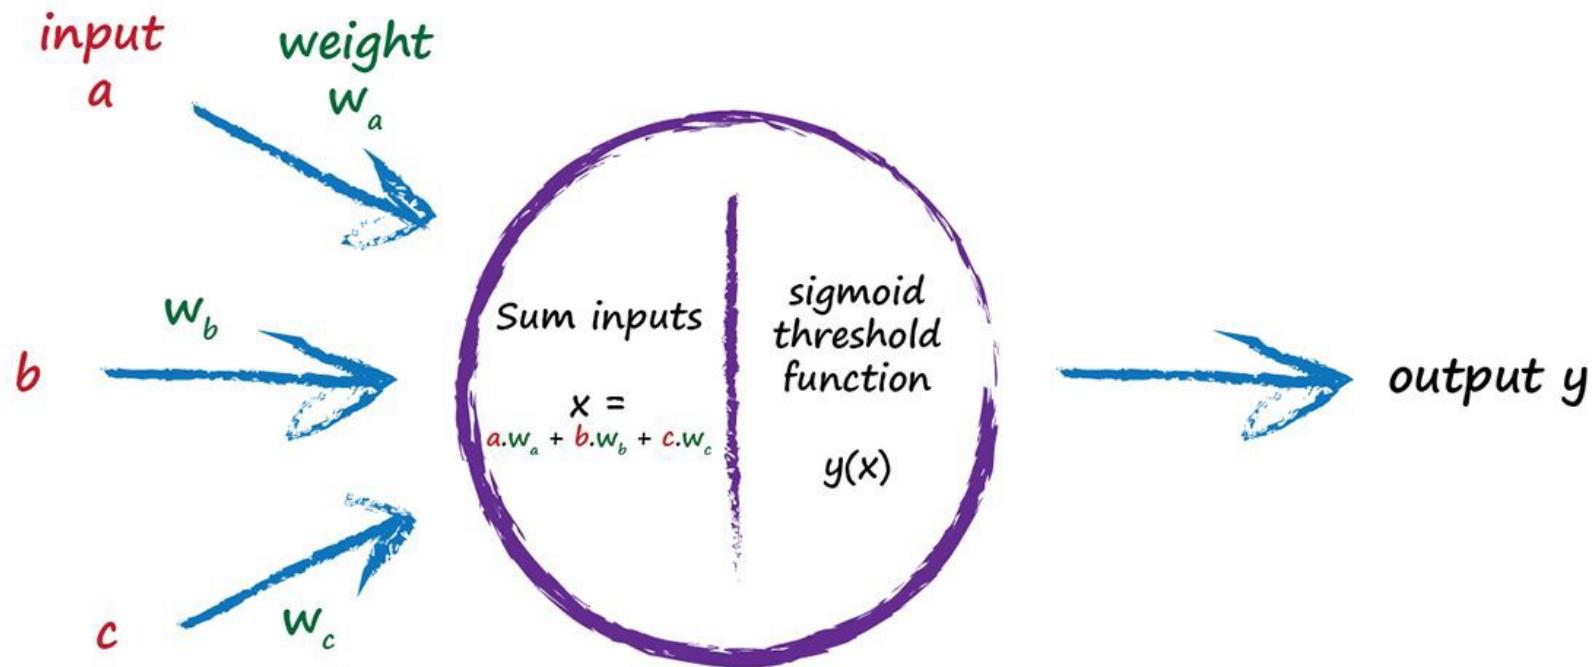


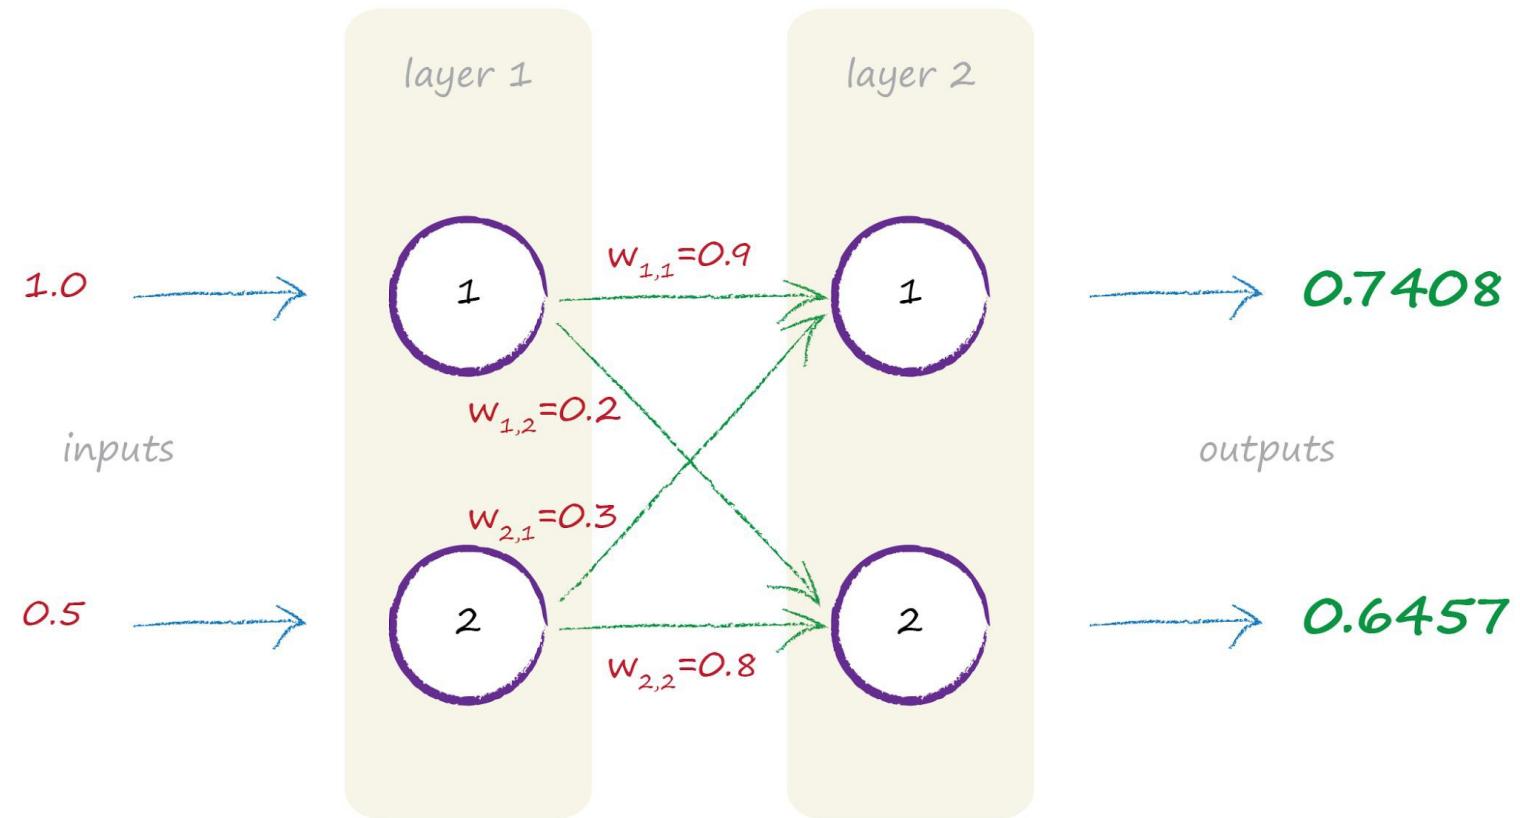


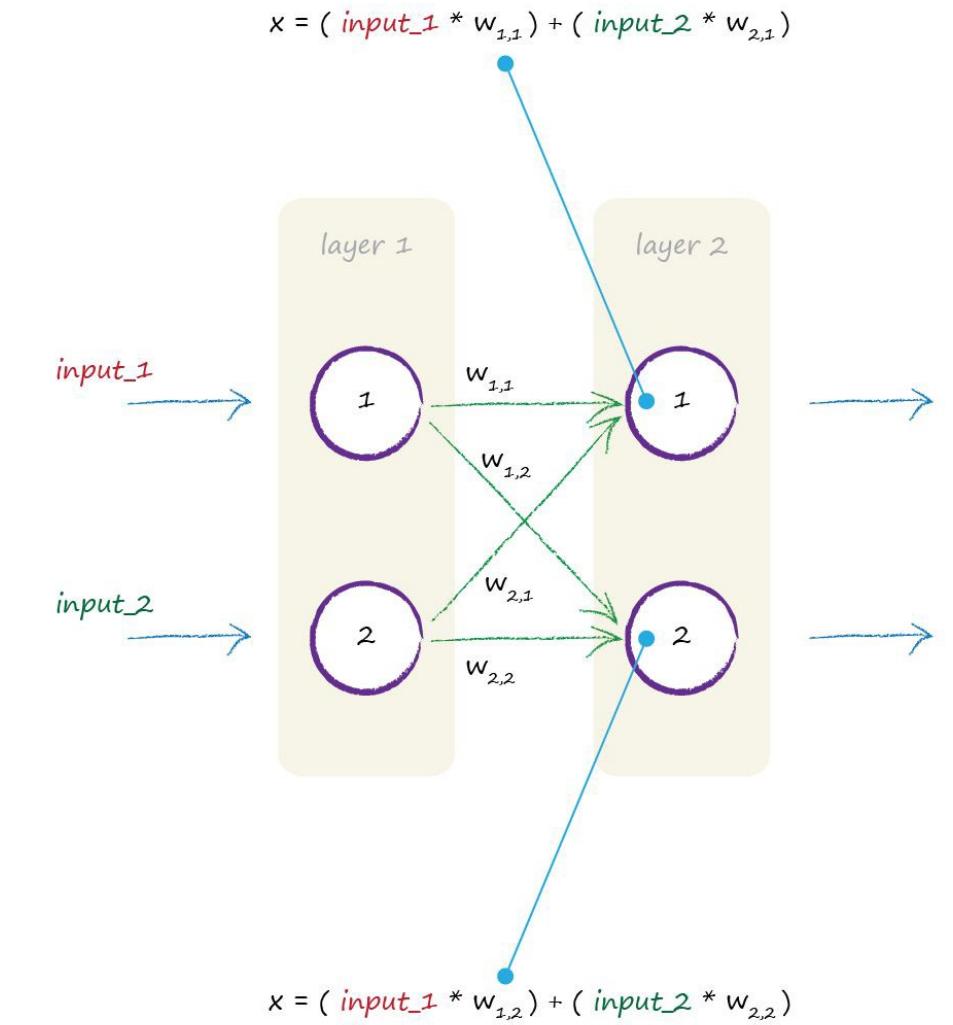


1. Natural brains can do **sophisticated** things, and are incredibly **resilient** to damage and imperfect signals .. unlike traditional computing.
2. Trying to copy biological brains partly inspired **artificial neural networks**.
3. **Link weights** are the adjustable parameter - it's where the learning happens.







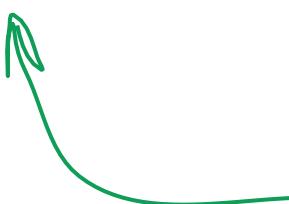


weights

incoming signals

$$\begin{pmatrix} w_{1,1} & w_{2,1} \\ w_{1,2} & w_{2,2} \end{pmatrix} \begin{pmatrix} \text{input_1} \\ \text{input_2} \end{pmatrix} = \begin{pmatrix} (\text{input_1} * w_{1,1}) + (\text{input_2} * w_{2,1}) \\ (\text{input_1} * w_{1,2}) + (\text{input_2} * w_{2,2}) \end{pmatrix}$$

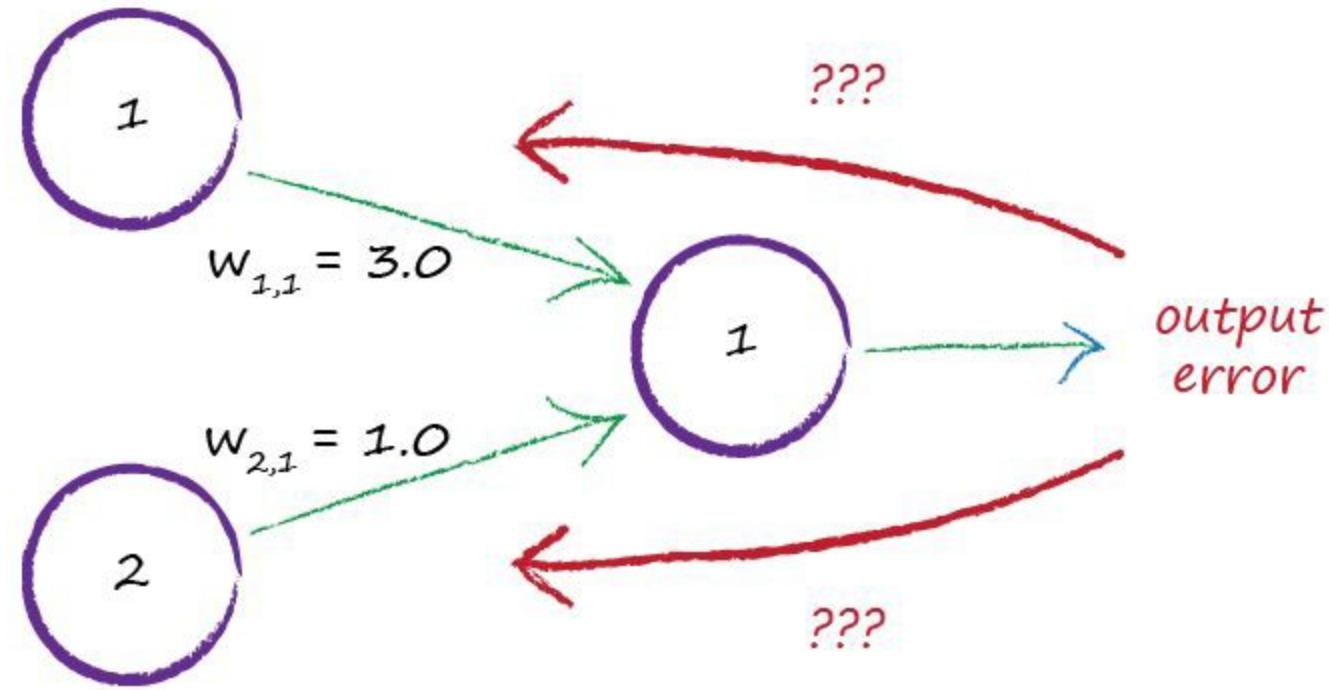
$$\mathbf{W} \cdot \mathbf{I} = \mathbf{X}$$

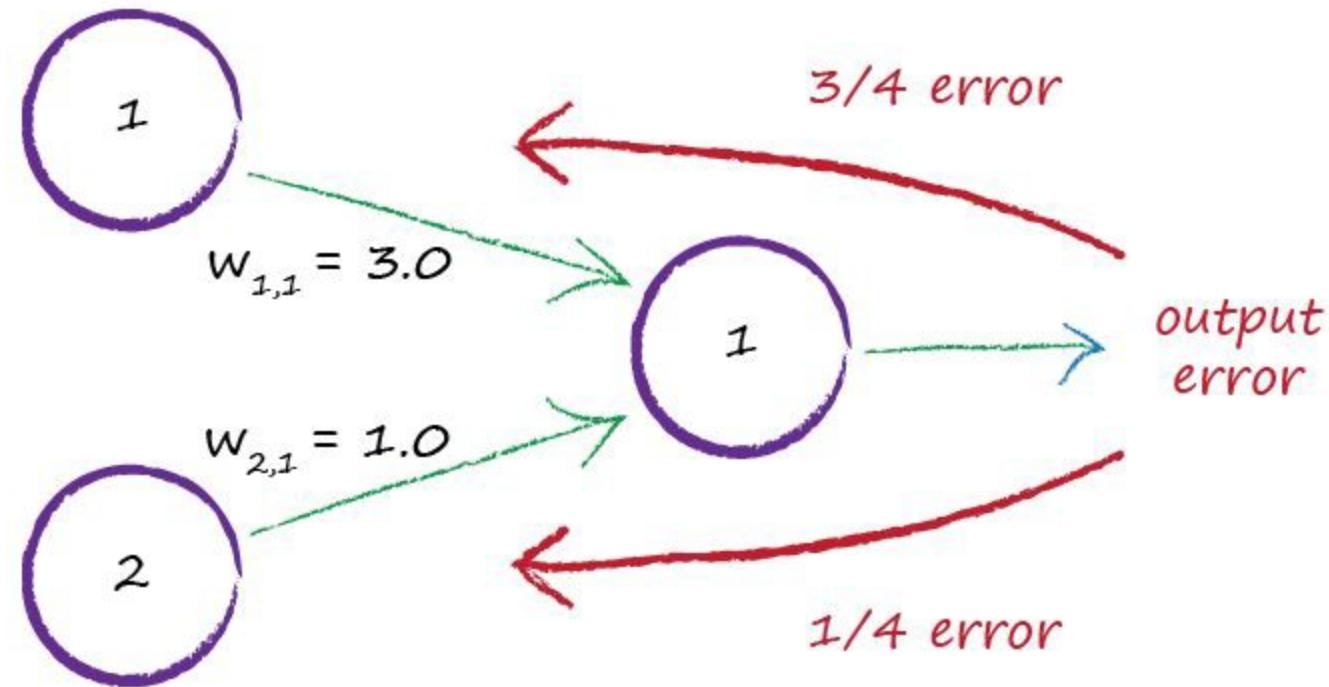


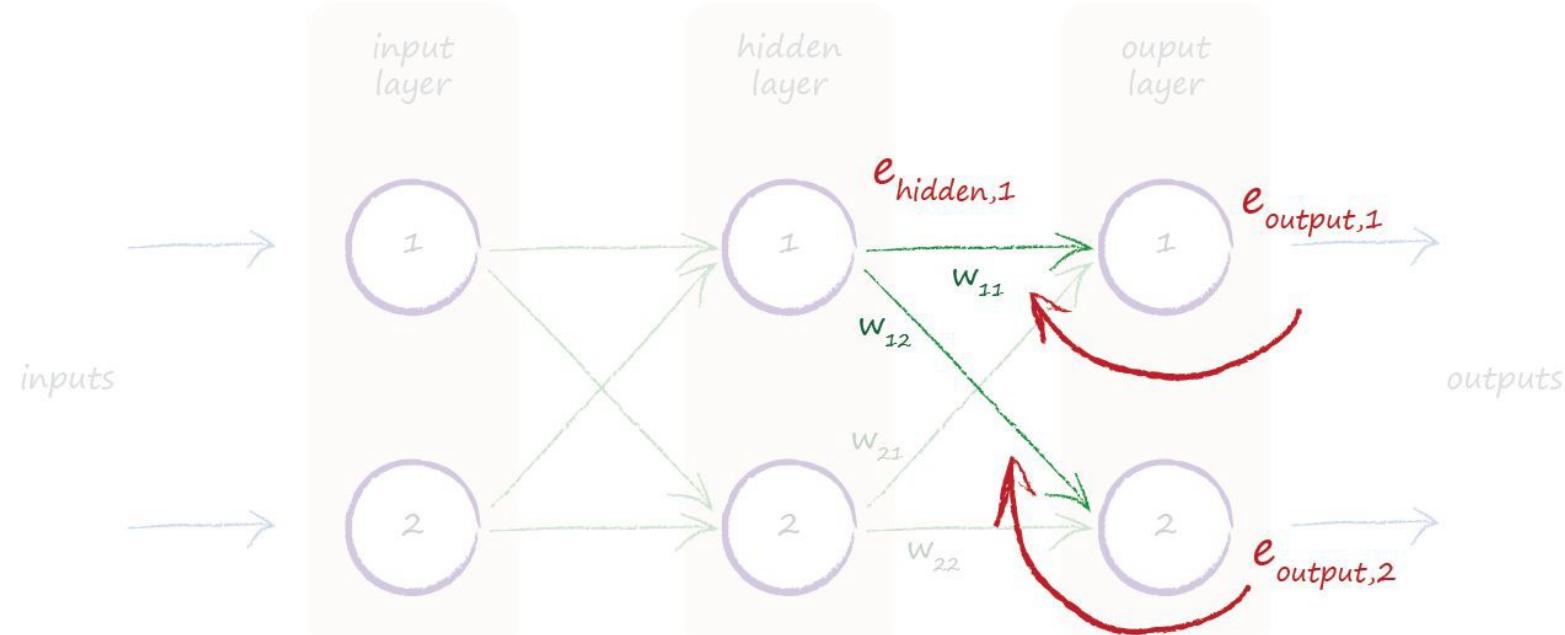
dot product

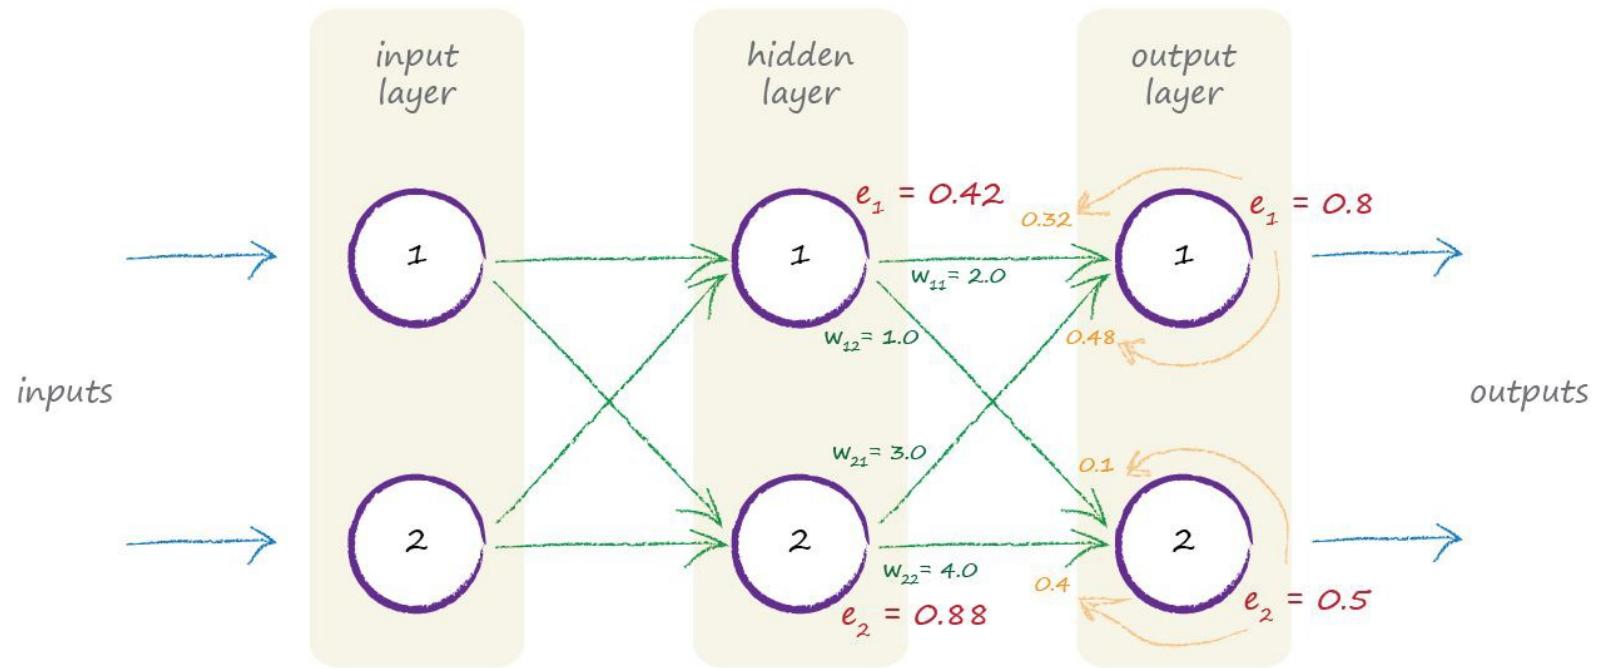


1. The many feedforward calculations can be expressed **concisely** as **matrix multiplication**, no matter what shape the network.
2. Some programming languages can do matrix multiplication really **efficiently** and **quickly**.







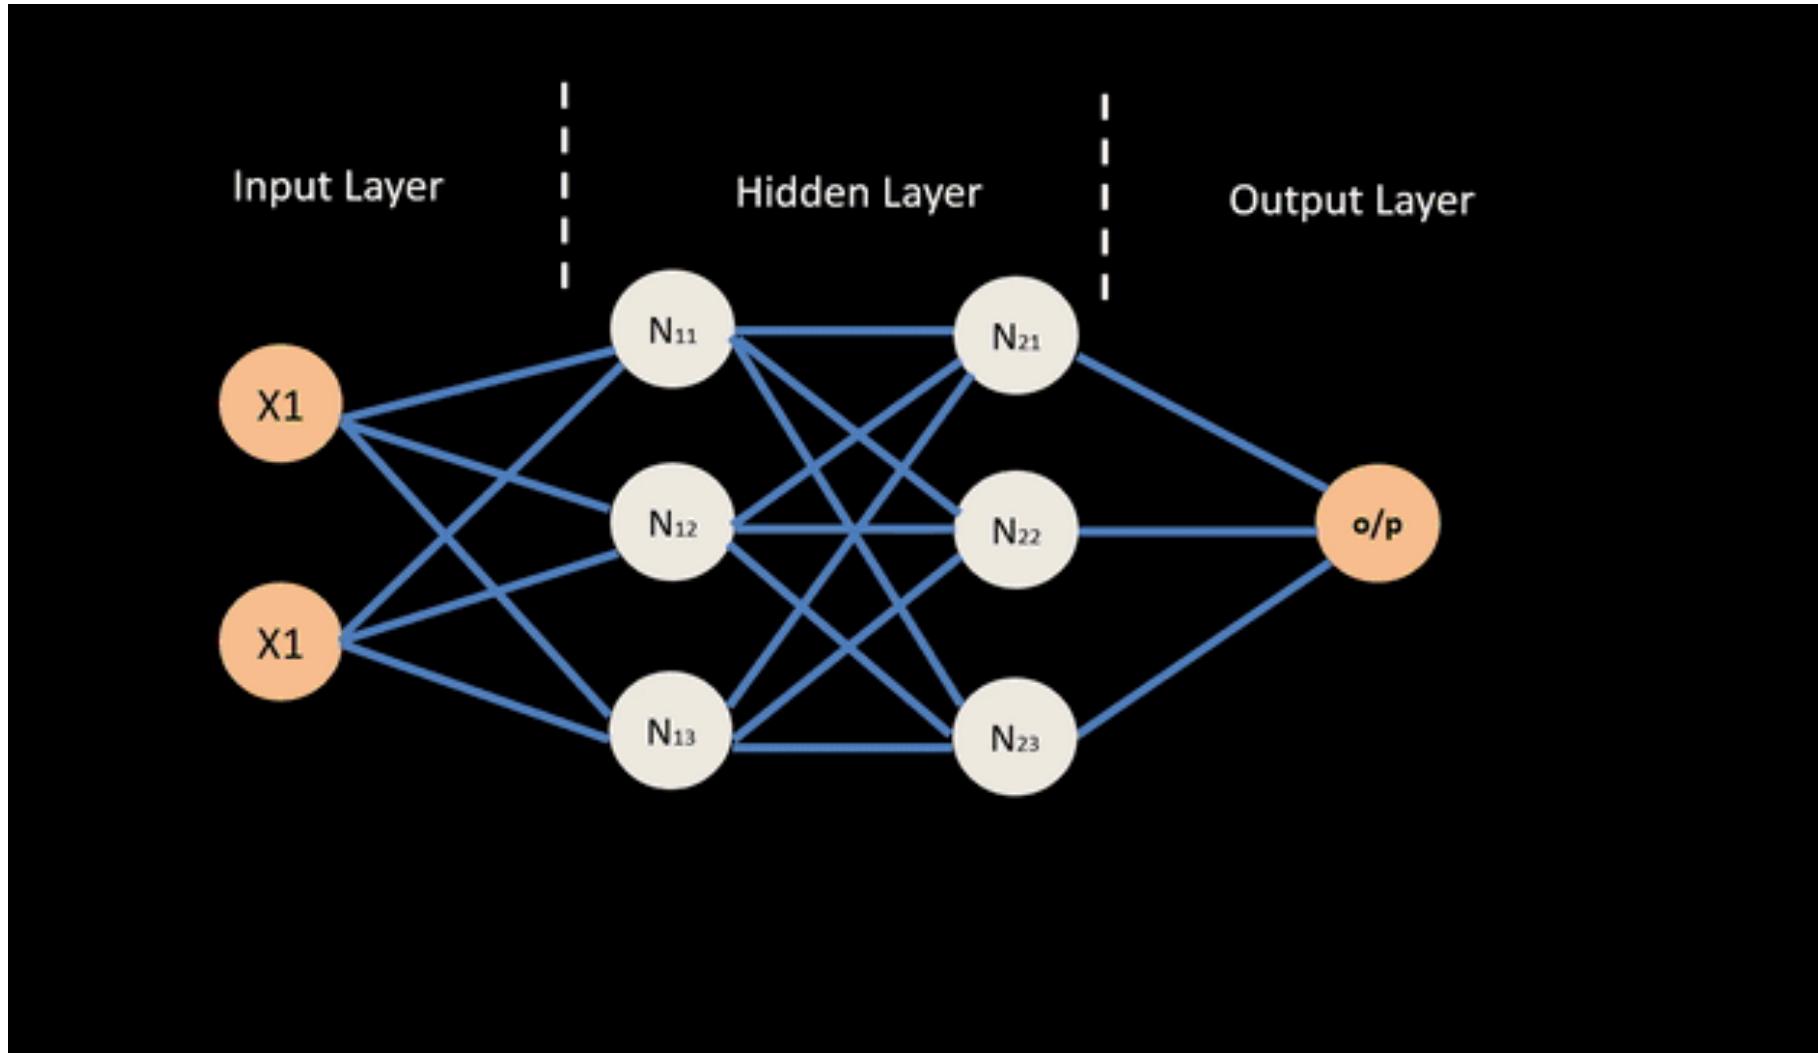


$$\text{error}_{\text{hidden}} = \begin{pmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{pmatrix} \cdot \begin{pmatrix} e_1 \\ e_2 \end{pmatrix}$$

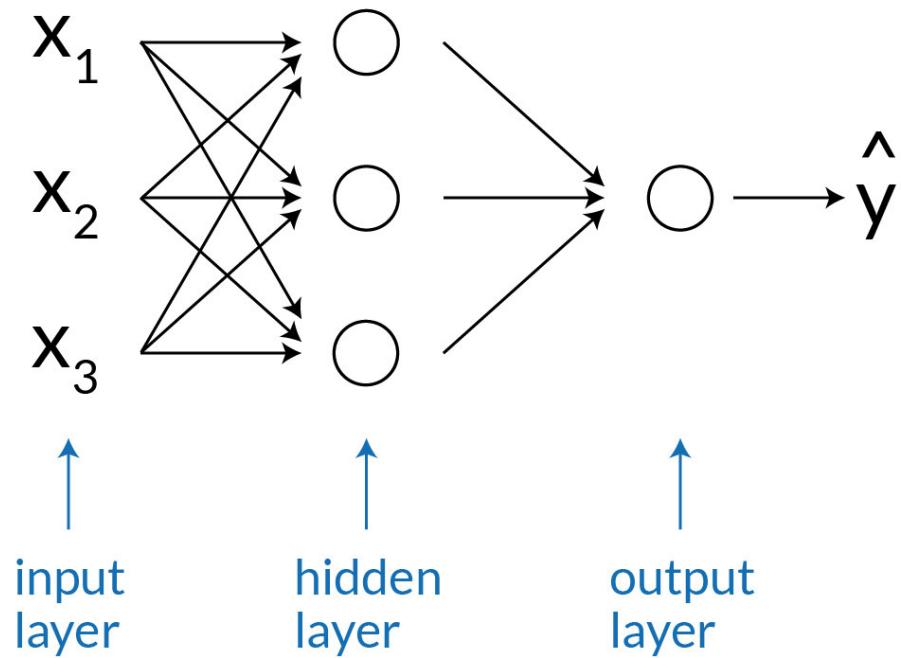
$$\text{error}_{\text{hidden}} = w^T_{\text{hidden_output}} \cdot \text{error}_{\text{output}}$$



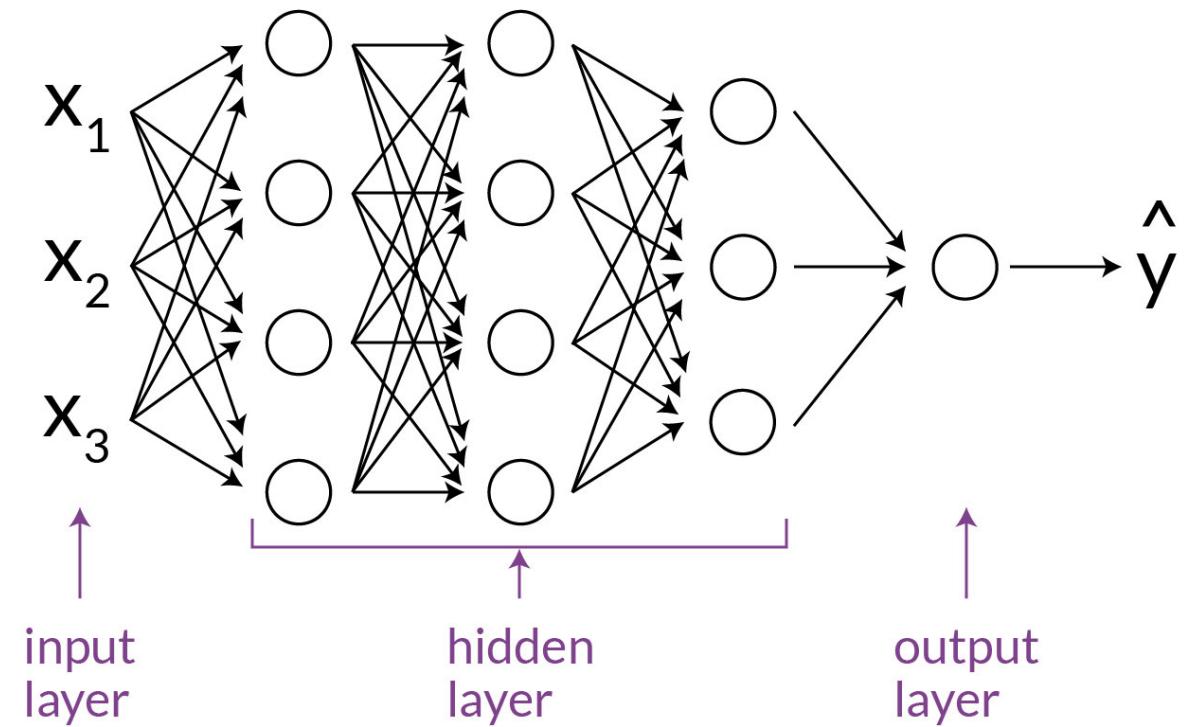
1. Remember we use the **error** to guide how we refine a model's parameter - link weights.
2. The error at the output nodes is easy - the difference between the **desired** and **actual** outputs.
3. The error at internal nodes isn't obvious. A **heuristic** approach is to split it in **proportion** to the link weights.
4. ... and back propagating the error can be expressed as a **matrix** multiplication too!



Shallow Neural Network

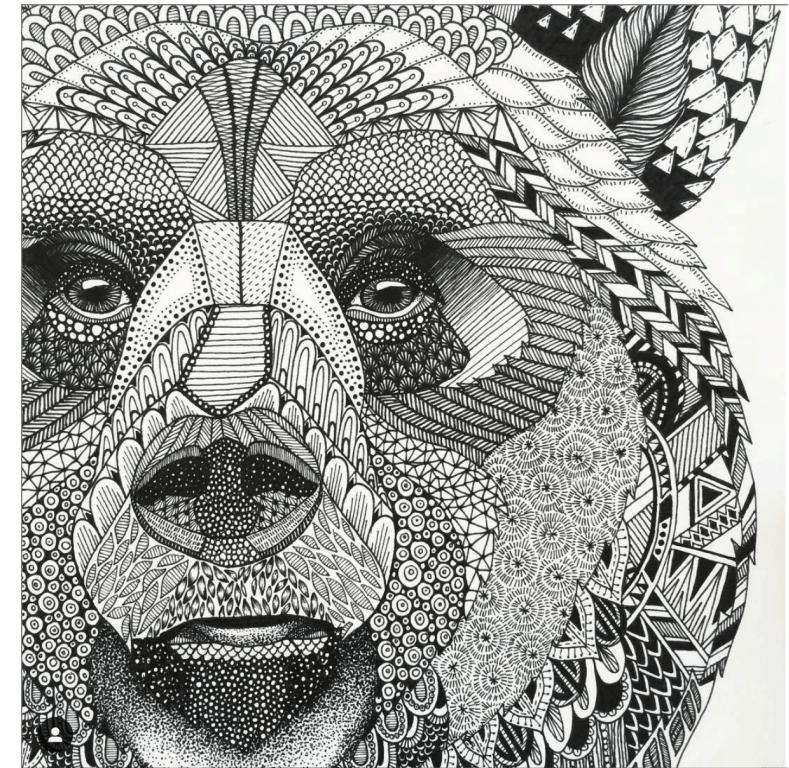


Deep Neural Network

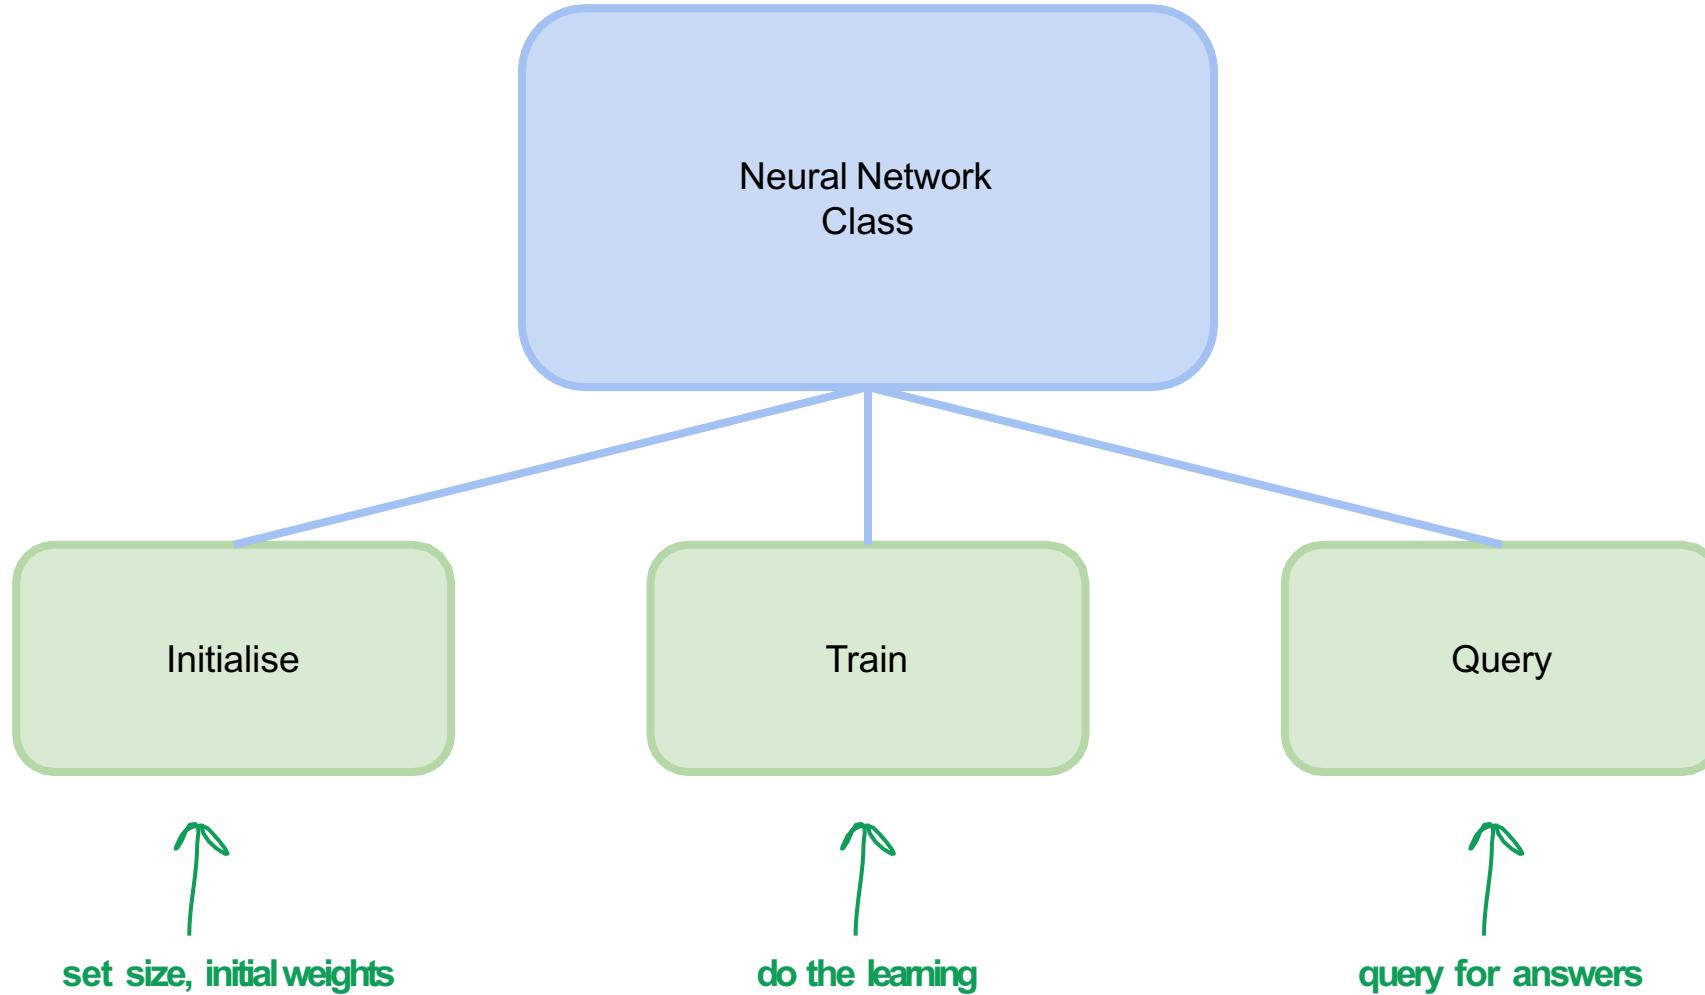


{04}

Do it Yourself



and Functions



numpy scipy
matplotlib
notebook

matrix maths



Function - Initialise

```
# initialise the neural network
def __init__(self, inputnodes, hiddennodes, outputnodes, learningrate): # set number
    of nodes in each input, hidden, output layer self.inodes = inputnodes
    self.hnodes = hiddennodes
    self.onodes = outputnodes

    # link weight matrices, wih and who
    # weights inside the arrays are w_i_j, where link is from node i to node j in the next layer # w11 w21
    # w12 w22 etc
    self.wih = numpy.random.normal(0.0, pow(self.hnodes, -0.5), (self.hnodes, self.inodes)) self.who =
    numpy.random.normal(0.0, pow(self.onodes, -0.5), (self.onodes, self.hnodes))

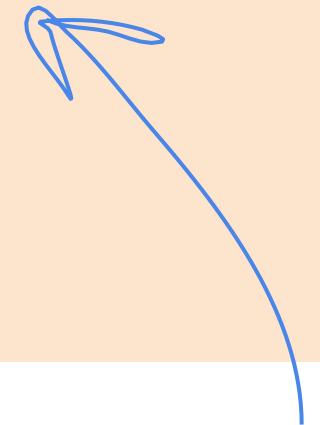
    # learning rate
    self.lr = learningrate

    # activation function is the sigmoid function
    self.activation_function = lambda x: scipy.special.expit(x)

    pass
```



numpy.random.normal()



random initial weights

Function - Query

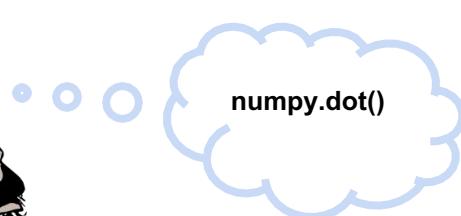
combined weighted signals into hidden layer
then sigmoid applied

```
# query the neural network
def query(self, inputs_list):
    # convert inputs list to 2d array
    inputs = numpy.array(inputs_list, ndmin=2).T

    # calculate signals into hidden layer
    hidden_inputs = numpy.dot(self.wih, inputs)
    # calculate the signals emerging from hidden layer
    hidden_outputs = self.activation_function(hidden_inputs)

    # calculate signals into final output layer
    final_inputs = numpy.dot(self.who, hidden_outputs)
    # calculate the signals emerging from final output layer
    final_outputs = self.activation_function(final_inputs)

    return final_outputs
```



similar for output layer

Function - Train

```
#train the neural network
def train(self, inputs_list, targets_list):
    #convert inputs list to 2d array
    inputs = numpy.array(inputs_list, ndmin=2).T
    targets = numpy.array(targets_list, ndmin=2).T
```

same feed forward as before

```
#calculate signals into hidden layer
hidden_inputs = numpy.dot(self.wih, inputs)
#calculate the signals emerging from hidden layer
hidden_outputs = self.activation_function(hidden_inputs)
```

output layer errors

```
#calculate signals into final output layer
final_inputs = numpy.dot(self.who, hidden_outputs)
#calculate the signals emerging from final output layer
final_outputs = self.activation_function(final_inputs)
```

hidden layer errors

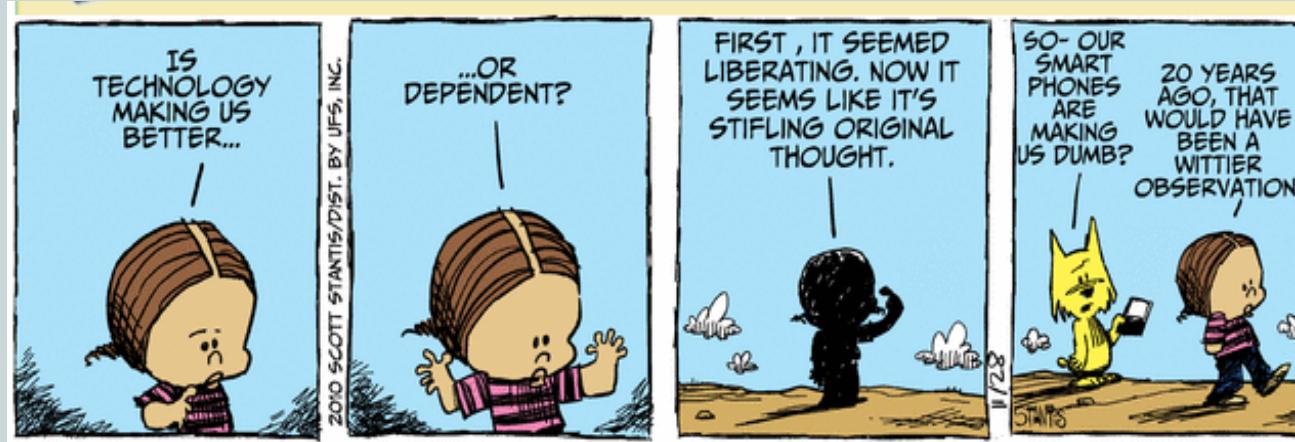
```
#output layer error is the (target - actual)
output_errors = targets - final_outputs
#hidden layer error is the output_errors, split by weights, recombined at hidden nodes
hidden_errors = numpy.dot(self.who.T, output_errors)
```

```
#update the weights for the links between the hidden and output layers
self.who += self.lr * numpy.dot((output_errors * final_outputs * (1.0 - final_outputs)),
numpy.transpose(hidden_outputs))
```

```
#update the weights for the links between the input and hidden layers
self.wih += self.lr * numpy.dot((hidden_errors * hidden_outputs * (1.0 - hidden_outputs)),
numpy.transpose(inputs))
```

pass

update weights



This lesson was developed by:

Robert Frans van der Willigen
CMD, Hogeschool Rotterdam
OKT 2020

<http://creativecommons.org/licenses/by-nc-sa/3.0/>

This lesson is licensed under a Creative Commons Attribution-Share-Alike license. You can change it, transmit it, show it to other people. Just always give credit to RFvdW.

Creative Commons License Types		
	Can someone use it commercially?	Can someone create new versions of it?
Attribution	ⓘ	ⓘ
Share Alike	ⓘ	Yup, AND they must license the new work under a Share Alike license.
No Derivatives	ⓘ	ⓘ
Non-Commercial	ⓘ	Yup, AND the new work must be non-commercial, but it can be under any non-commercial license.
Non-Commercial Share Alike	ⓘ	Yup, AND they must license the new work under a Non-Commercial Share Alike license.
Non-Commercial No Derivatives	ⓘ	ⓘ

SOURCE
<http://www.masternewmedia.org/how-to-publish-a-book-under-a-creative-commons-license/>

<http://empoweringthenatives.edublogs.org/2012/03/15/creative-commons-licenses/>

<http://creativecommons.org/licenses/>

